



Supervised Learning of Neural Random-Access Machines with Differential Evolution

Valerio Belli

April 11, 2018

University of Perugia

Table of contents

1. Neural Random-Access Machines
2. Artificial Neural Network
3. Differential Evolution & DENN
4. Implementation & results

Neural Random-Access Machines

What is?

It is a machine introduced in [Kurach et al., 2015] based on a neural network which is capable of manipulating pointers and dereferencing them through “logical” circuits. Its objective is to solve a task on which it has been trained creating and executing those circuits.

Introduction

Let $N = \{1, \dots, l - 1\}$, where l is an integer constant, the integers set over the NRAM should work. Since the training in [Kurach et al., 2015] is made through a gradient descent algorithm, the NRAM does not work directly over N but over stochastically independent probability distributions, defined as $p \in \mathbb{R}^l$ satisfying $0 \leq p_i \leq 1$ and $\sum_{i=0}^{l-1} p_i = 1$. Moreover, let T the maximum number of timestep of execution of the NRAM.

Registers only version

Controller

Registers

Modules

Zero

One

Two

Inc

Add

Sub

Dec

Less-
Than

Less-
Equal-Than

Equal-
Than

Min

Max

Modules

Controller

Registers

Modules

Zero

One

Two

Inc

Add

Sub

Dec

Less-
Than

Less-
Equal-Than

Equal-
Than

Min

Max

Modules

The modules (or gates) are components through which the controller, connecting them, manipulates values and pointers. In the NRAM exist three types of modules, defined as follows:

$$m_i \in N \text{ (Constant modules)} \quad (1)$$

$$m_i : p \rightarrow p \text{ (Unary modules)} \quad (2)$$

$$m_i : p \times p \rightarrow p \text{ (Binary modules)} \quad (3)$$

Each of them takes as input and returns probability distributions, as exception are constant modules which return only probability distributions.

Registers

Controller

Registers

Modules

Zero

One

Two

Inc

Add

Sub

Dec

Less-
Than

Less-
Equal-Than

Equal-
Than

Min

Max

Registers

The registers are a set of memory cells, where each of them contains a probability distribution. Hence, in other words, every register play the role of a random variable.

Controller

Controller

Registers

Modules

Zero

One

Two

Inc

Add

Sub

Dec

Less-
Than

Less-
Equal-Than

Equal-
Than

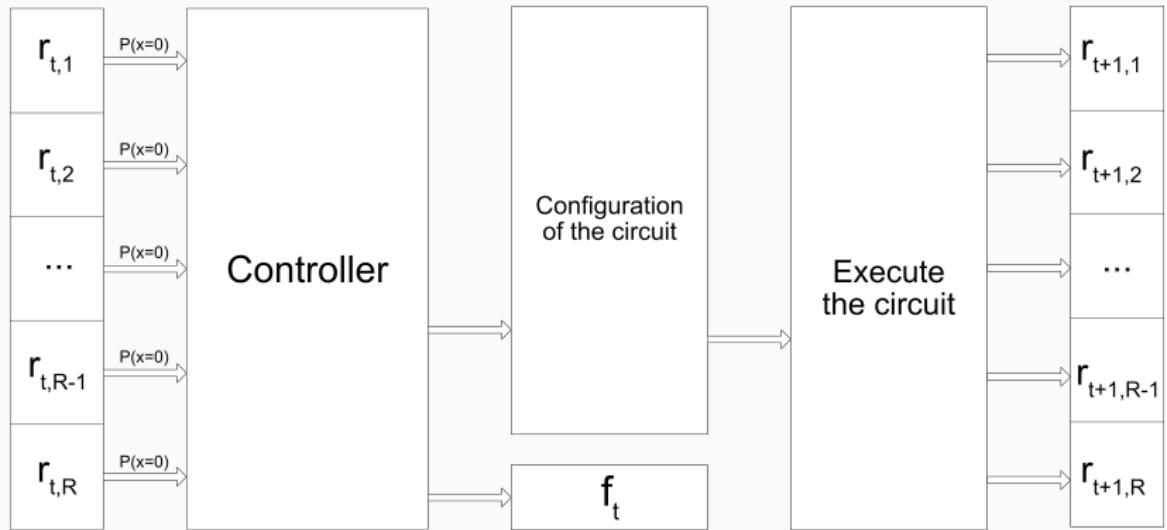
Min

Max

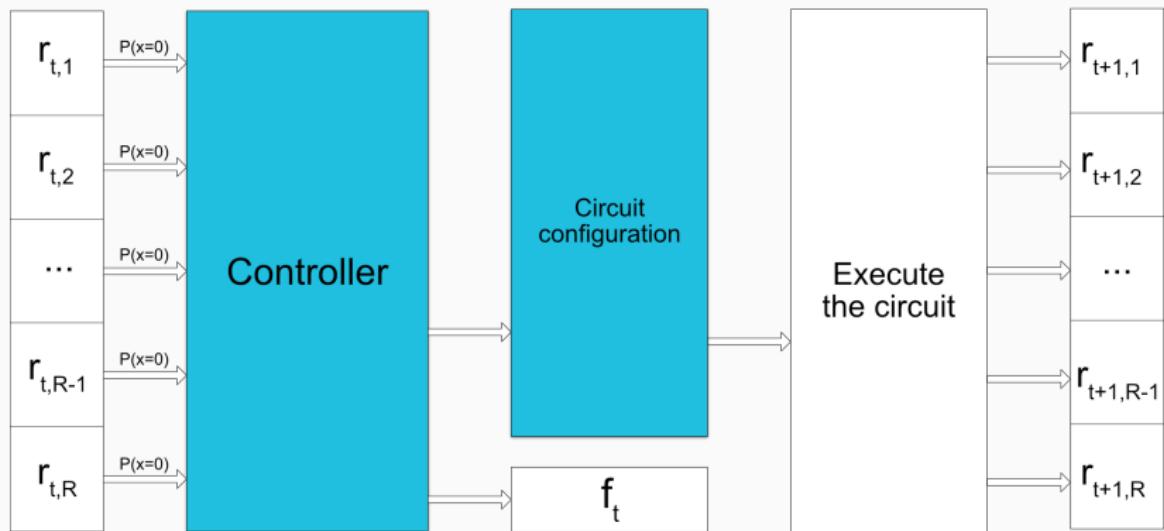
Controller

The controller is a neural network (Multi Layer Perceptron) whose objective is to emit a circuit configuration, i.e. how the circuits are connected together in form of a Directed Acyclic Graph. To do this the controller takes as input the $\mathbb{P}(r_i = 0)$ of each register.

Timestep execution

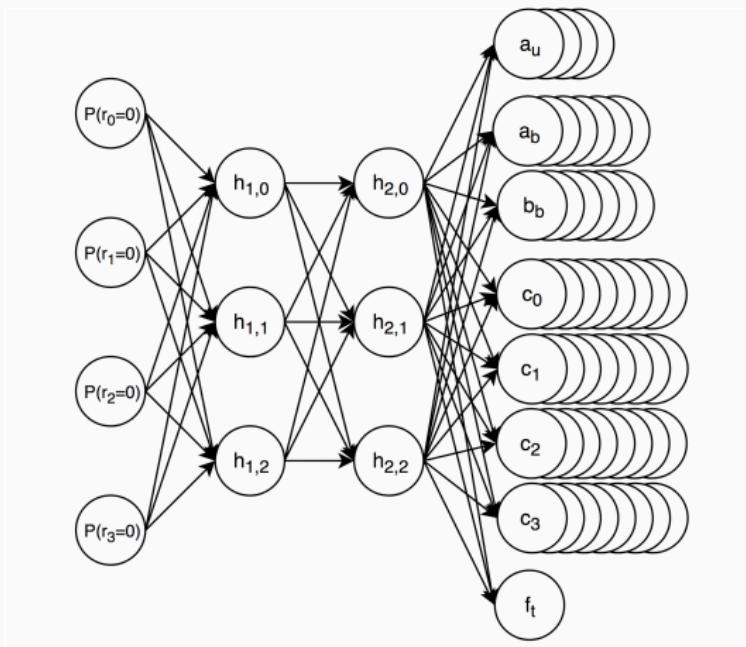


Timestep execution - Controller & Circuit configuration structure



Timestep execution - Controller

Let $R = 4$ and Gates= $\{m_u, m_c, m_b\}$, the controller structure appear as follows



Timestep execution - Circuit configuration structure

Using the same assumptions seen in the previous slide, the structure of the circuit configuration is represented as a list of vectors as follows

| | | | | | | | |
|-----------------------|-----|-----|-----|-----|-----|-----|---|
| $\text{softmax}(a_u)$ | 0 | 0 | 0 | 1 | | | |
| $\text{softmax}(a_b)$ | 0 | 0 | 0.4 | 0 | 0.7 | 0 | |
| $\text{softmax}(b_b)$ | 0 | 0.4 | 0.1 | 0.1 | 0.3 | 0.1 | |
| $\text{softmax}(c_0)$ | 0.2 | 0.2 | 0.3 | 0.1 | 0.2 | 0 | 0 |
| $\text{softmax}(c_1)$ | 0.3 | 0.1 | 0.1 | 0.4 | 0 | 0.1 | 0 |
| $\text{softmax}(c_2)$ | 0.6 | 0.3 | 0.1 | 0 | 0 | 0 | 0 |
| $\text{softmax}(c_3)$ | 0 | 1 | 0 | 0 | 0 | 0 | 0 |
| $\sigma(f_t)$ | 0.3 | | | | | | |

Timestep execution - Selection of the input of the gates

An input for the gate m_i is chosen by the controller from the set $\{r_1, \dots, r_R, o_1, \dots, o_{i-1}\}$ where:

- r_j is the content of the j^{th} register, $j = 1, \dots, R$;
- o_k is the output of the k^{th} previous gate with respect to the current m_i , for $k = 1, \dots, |Gates| - 1$.

selected through a weighted average with a coefficient s_i as follows

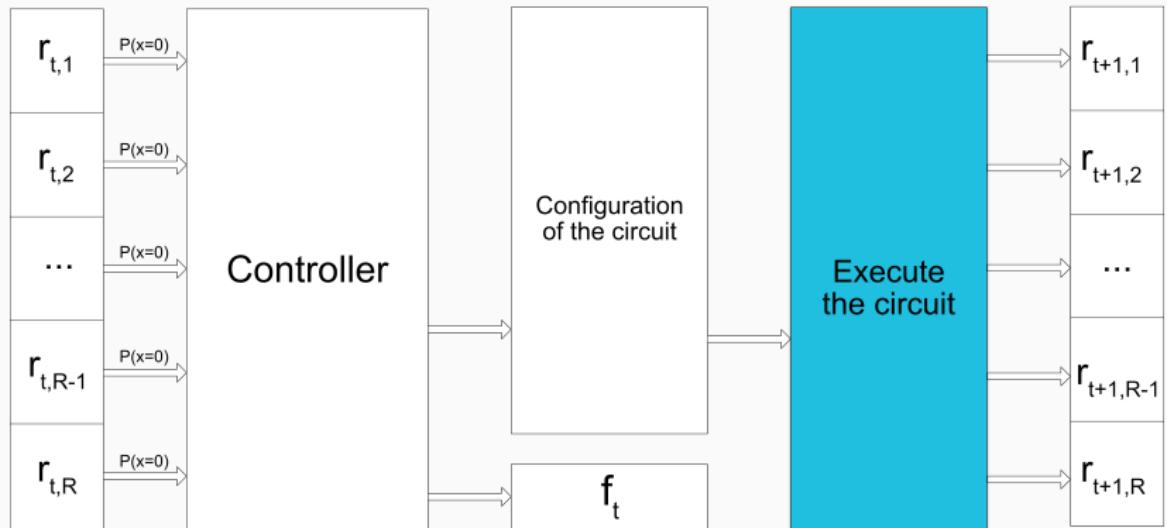
$$(r_1, r_2, \dots, r_R, o_1, o_2, \dots, o_{i-1})^T \text{softmax}(s_i)$$

Timestep execution - Selection of the input of the registers

Similarly to the gates, also the content of the registers is updated as follows:

$$r_i^{(t+1)} = (r_1^{(t)}, \dots, r_R^{(t)}, o_1^{(t)}, \dots, o_{|\text{Gates}|}^{(t)})^T \text{softmax}(s_i)$$

Timestep execution - Circuit execution example



Timestep execution - Circuit execution example



Timestep execution - Circuit execution example

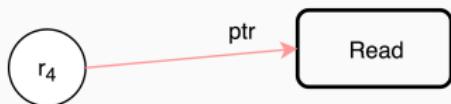
r_4

Read

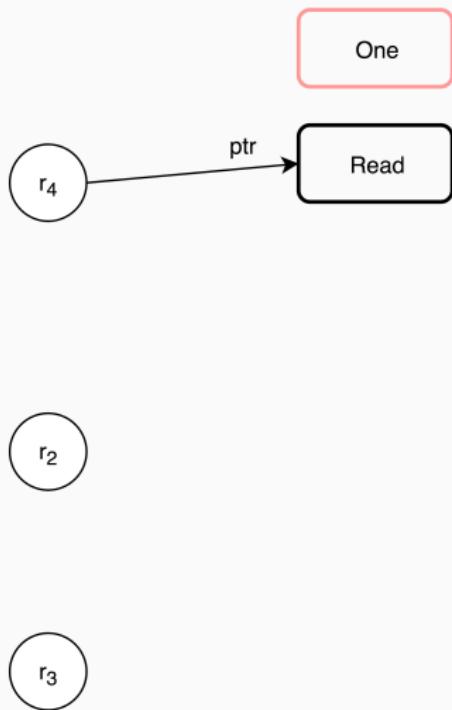
r_2

r_3

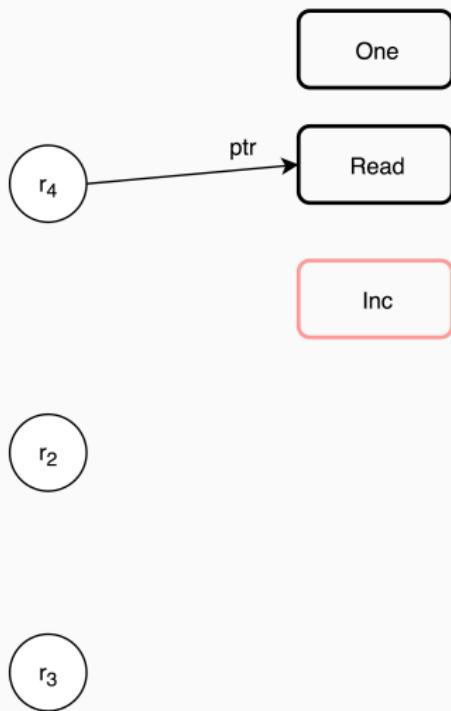
Timestep execution - Circuit execution example



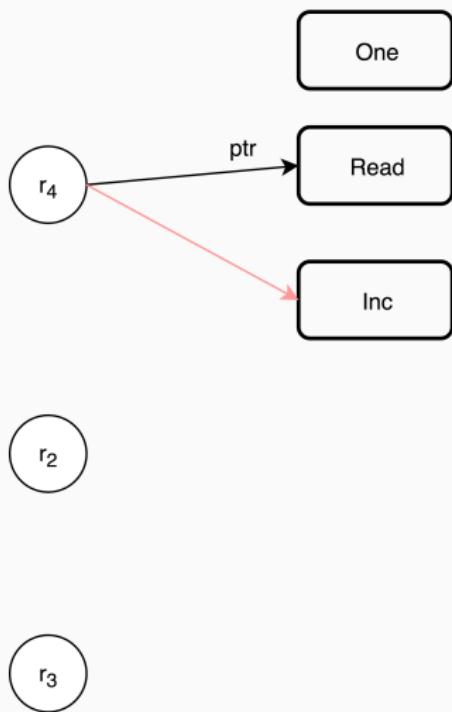
Timestep execution - Circuit execution example



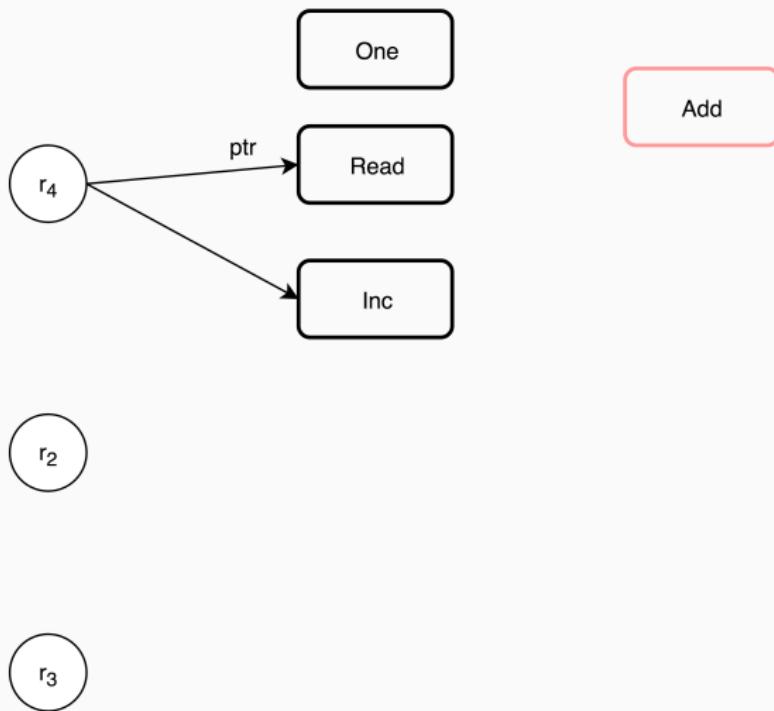
Timestep execution - Circuit execution example



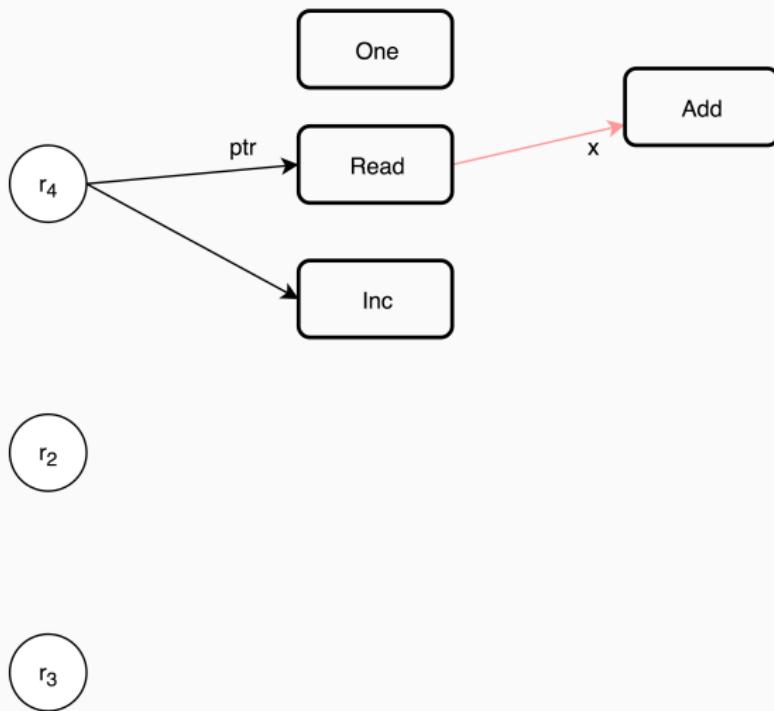
Timestep execution - Circuit execution example



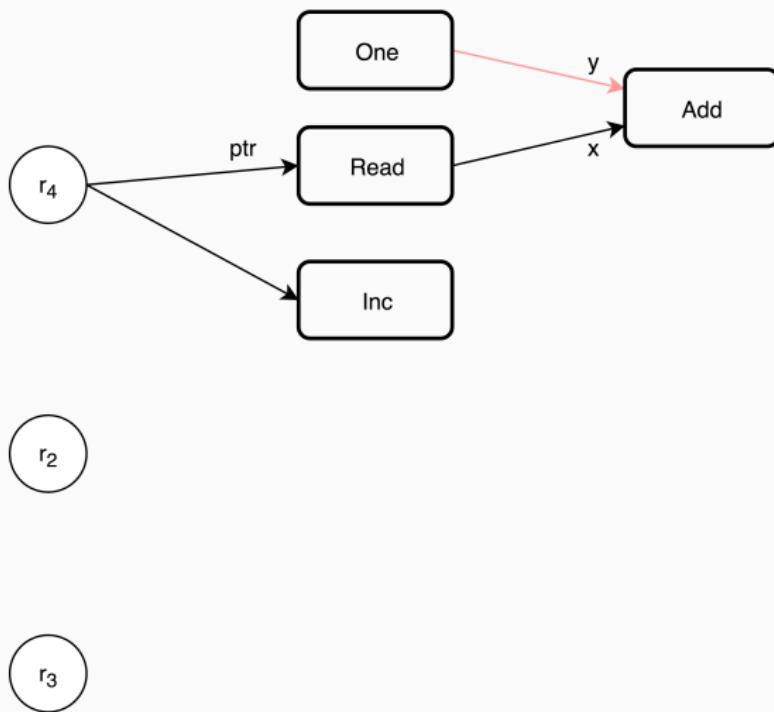
Timestep execution - Circuit execution example



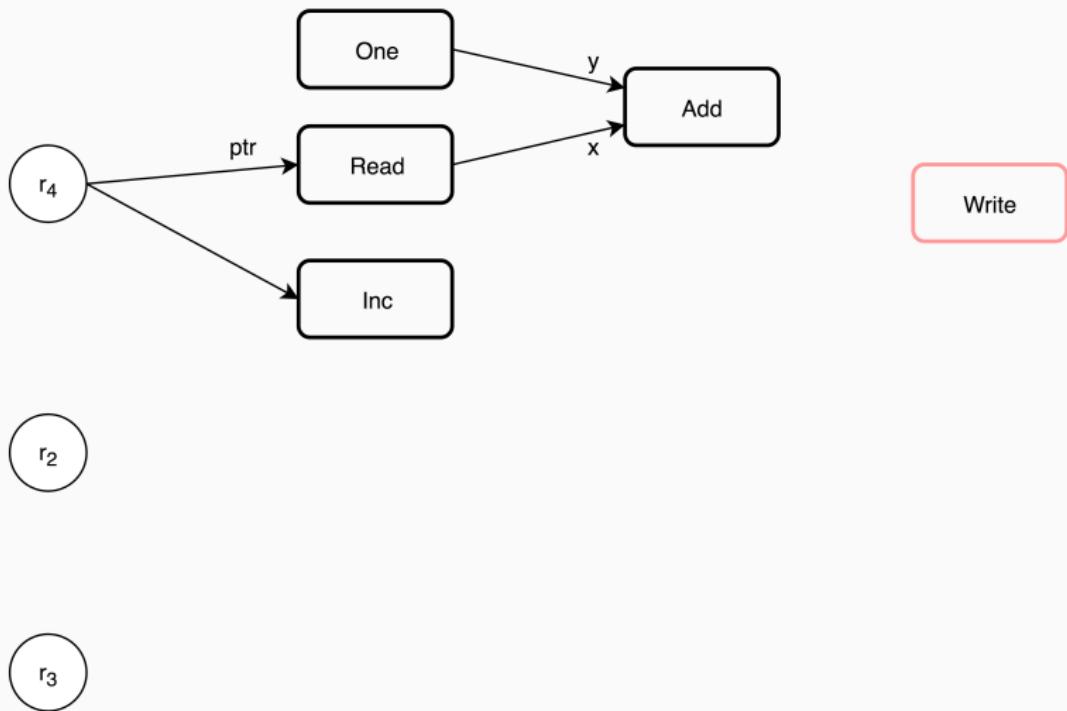
Timestep execution - Circuit execution example



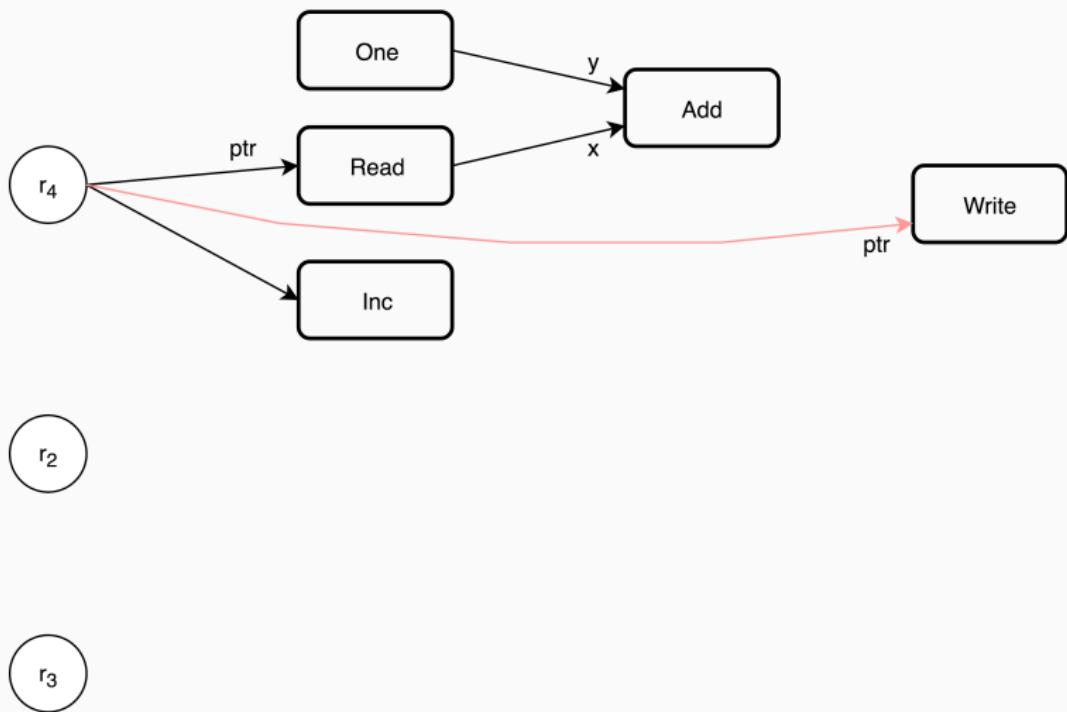
Timestep execution - Circuit execution example



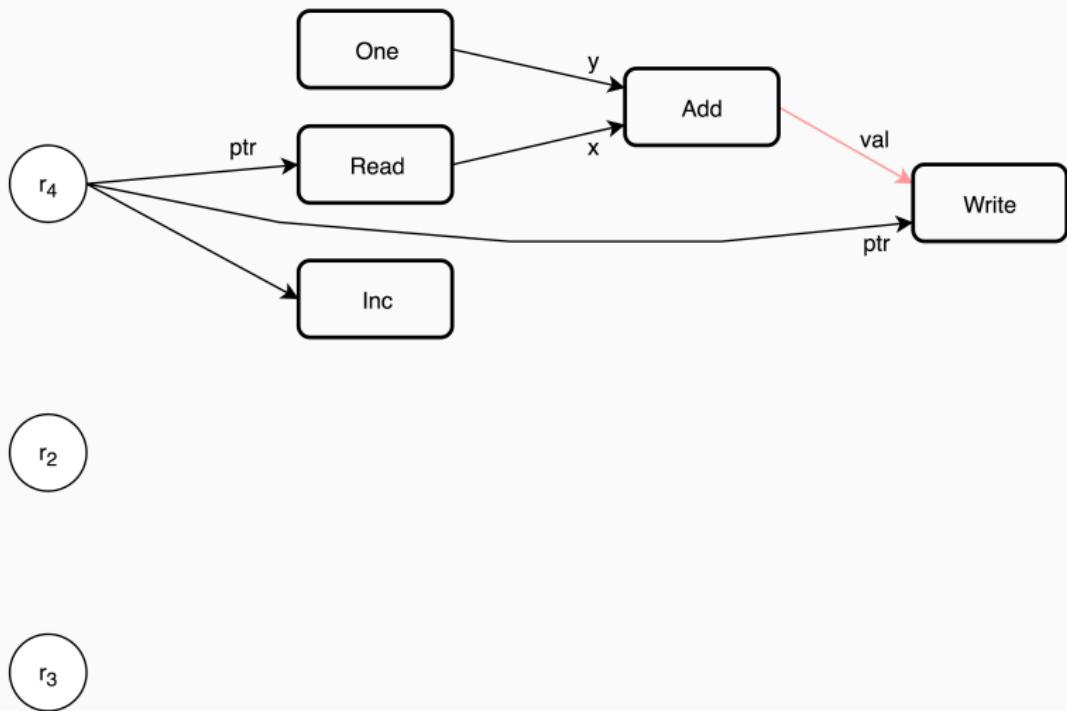
Timestep execution - Circuit execution example



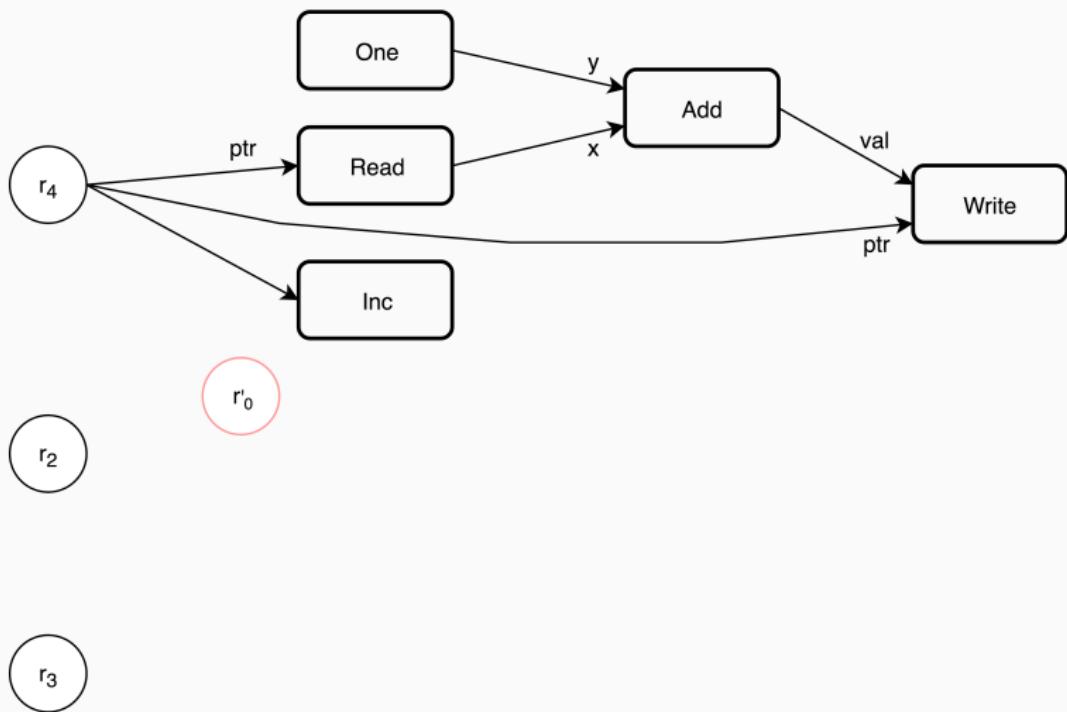
Timestep execution - Circuit execution example



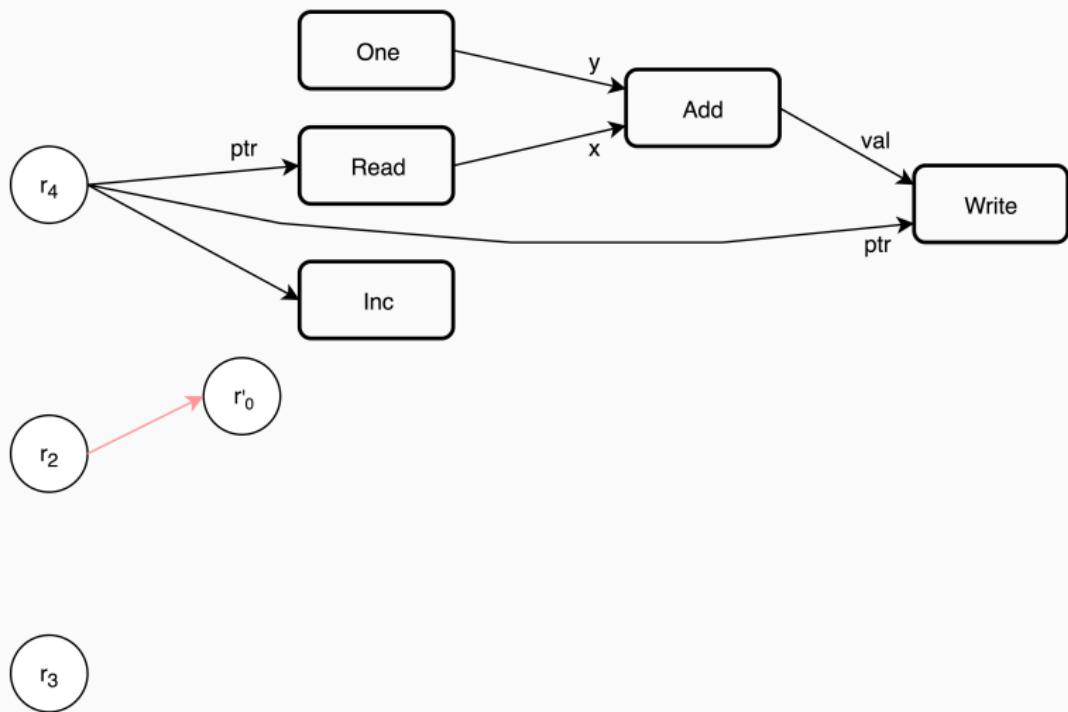
Timestep execution - Circuit execution example



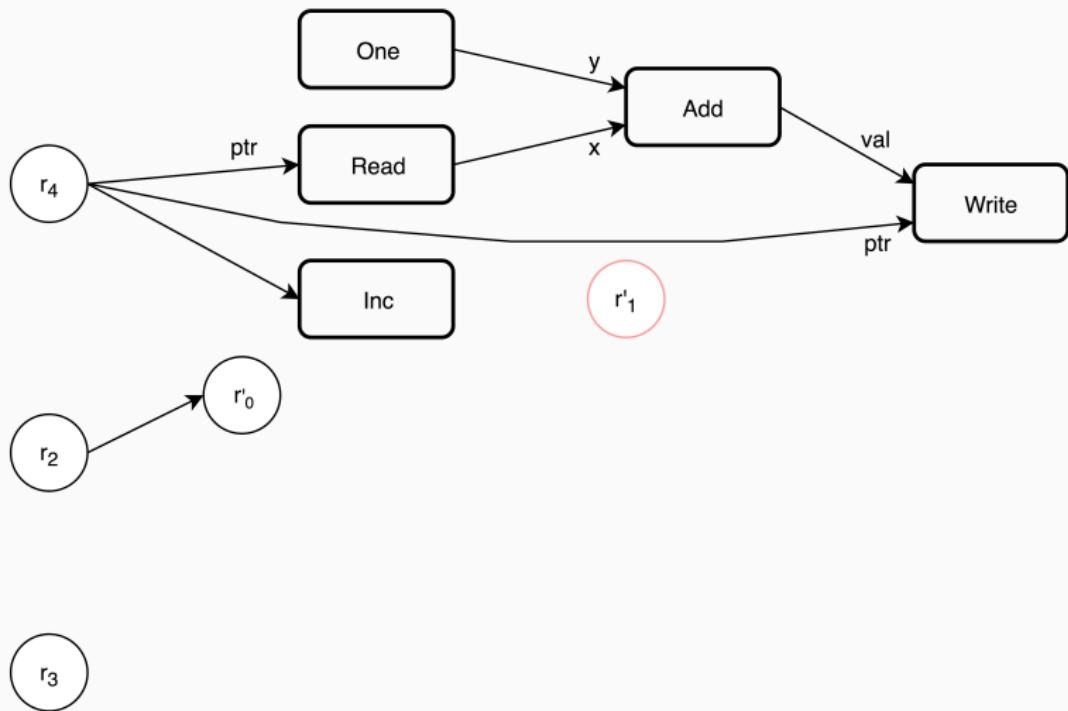
Timestep execution - Circuit execution example



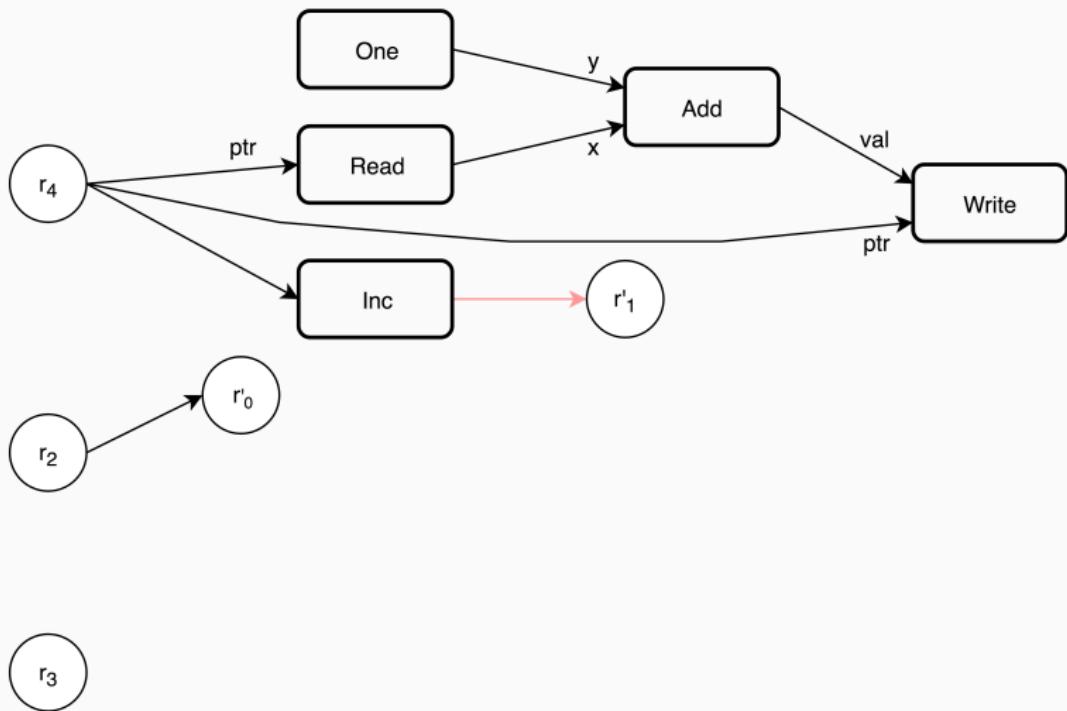
Timestep execution - Circuit execution example



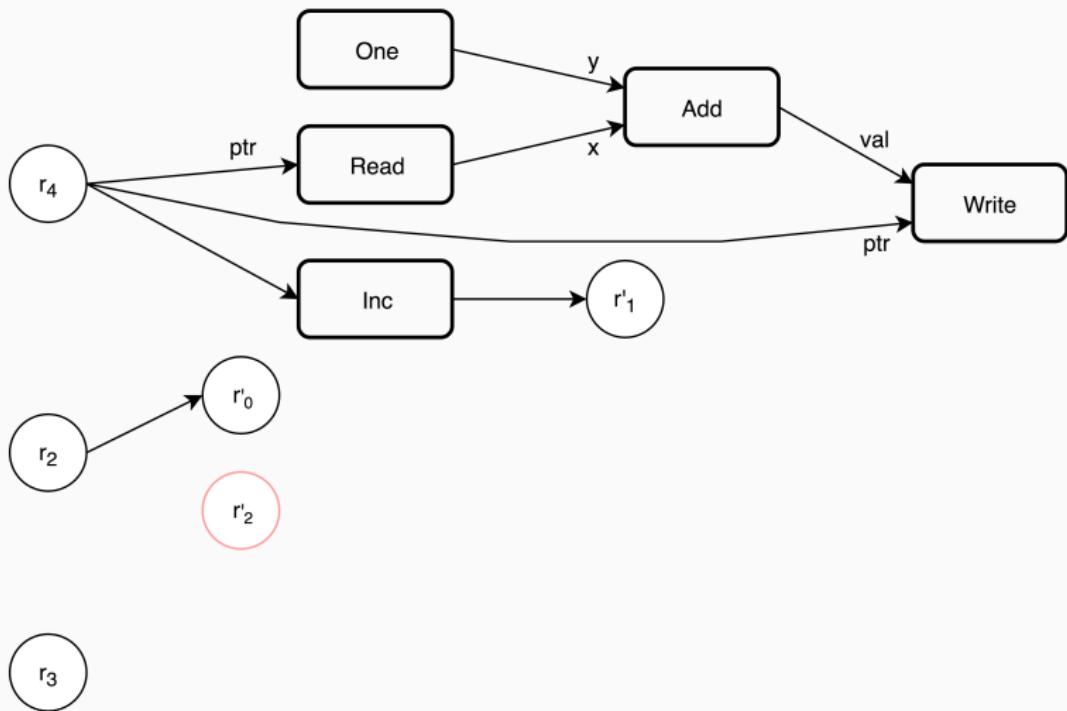
Timestep execution - Circuit execution example



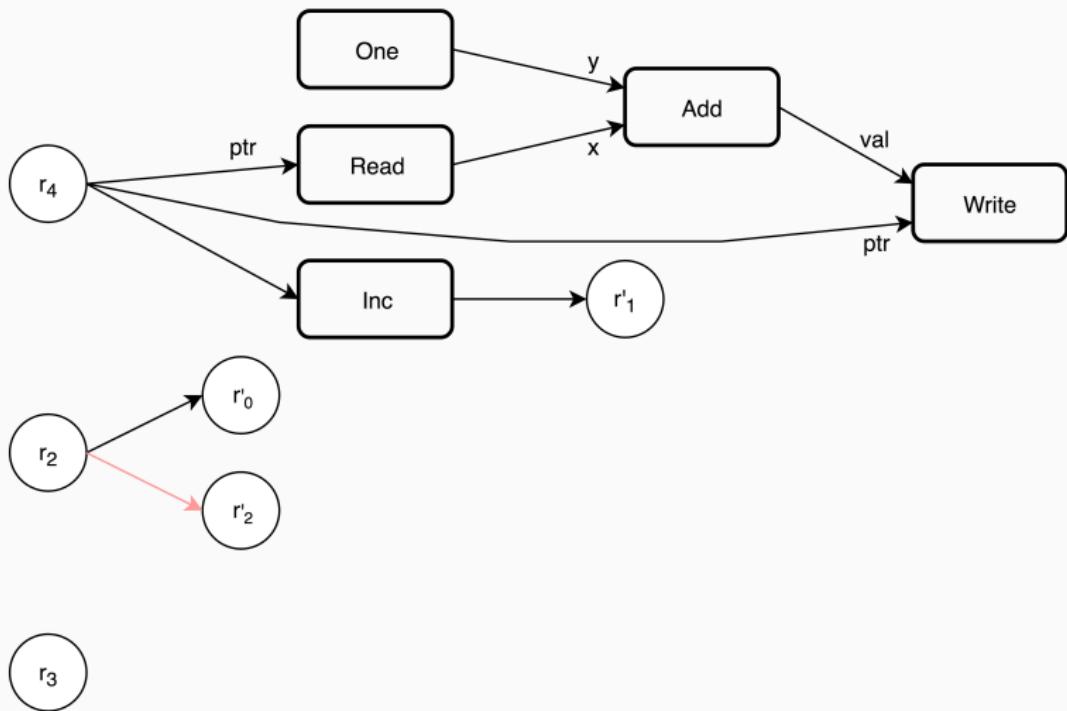
Timestep execution - Circuit execution example



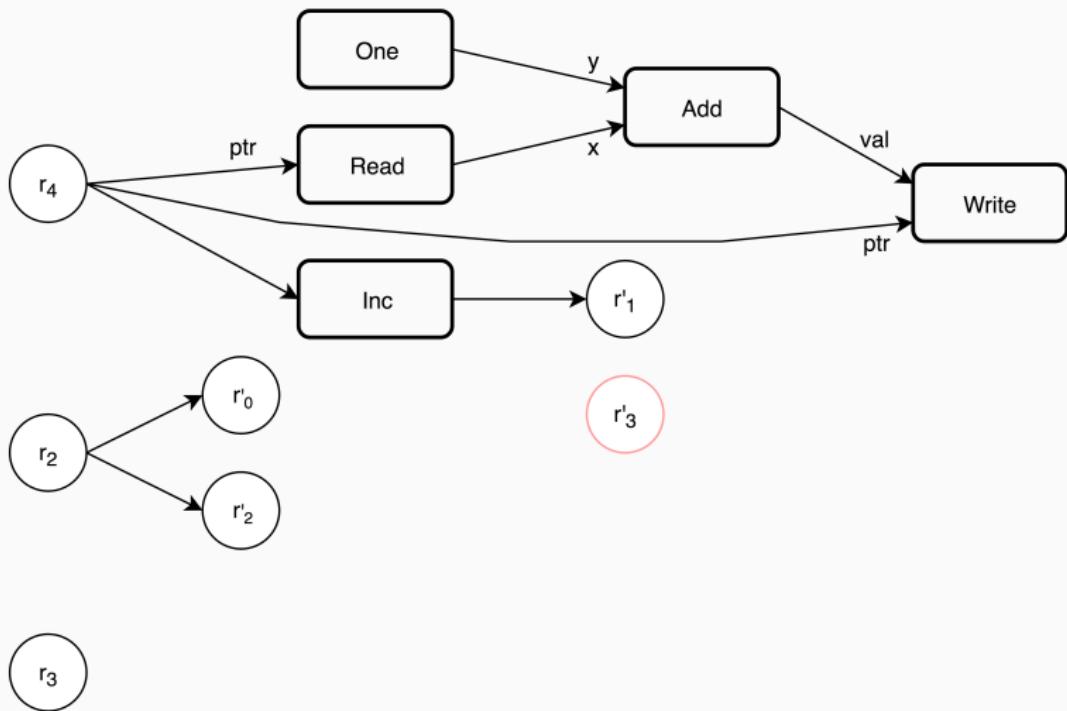
Timestep execution - Circuit execution example



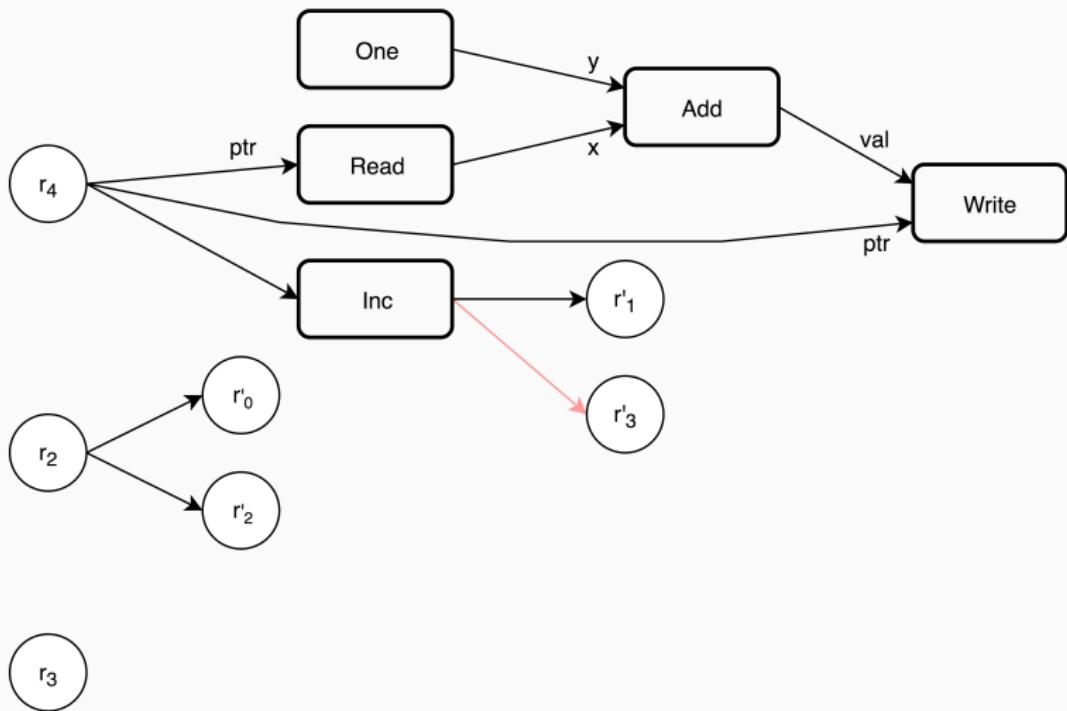
Timestep execution - Circuit execution example



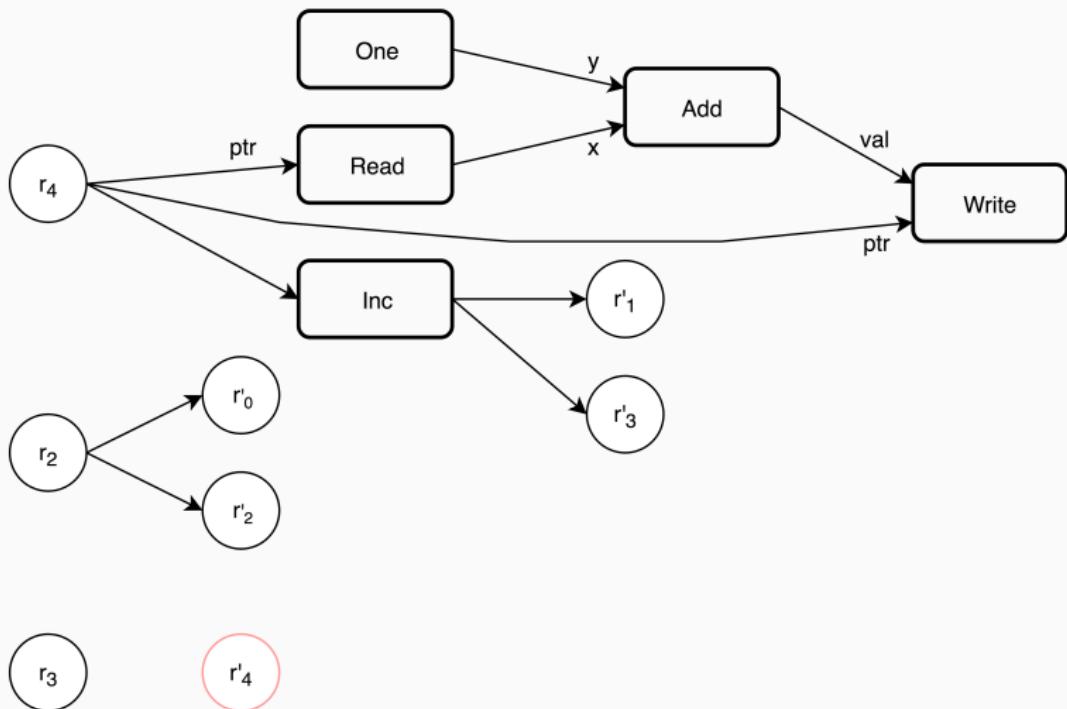
Timestep execution - Circuit execution example



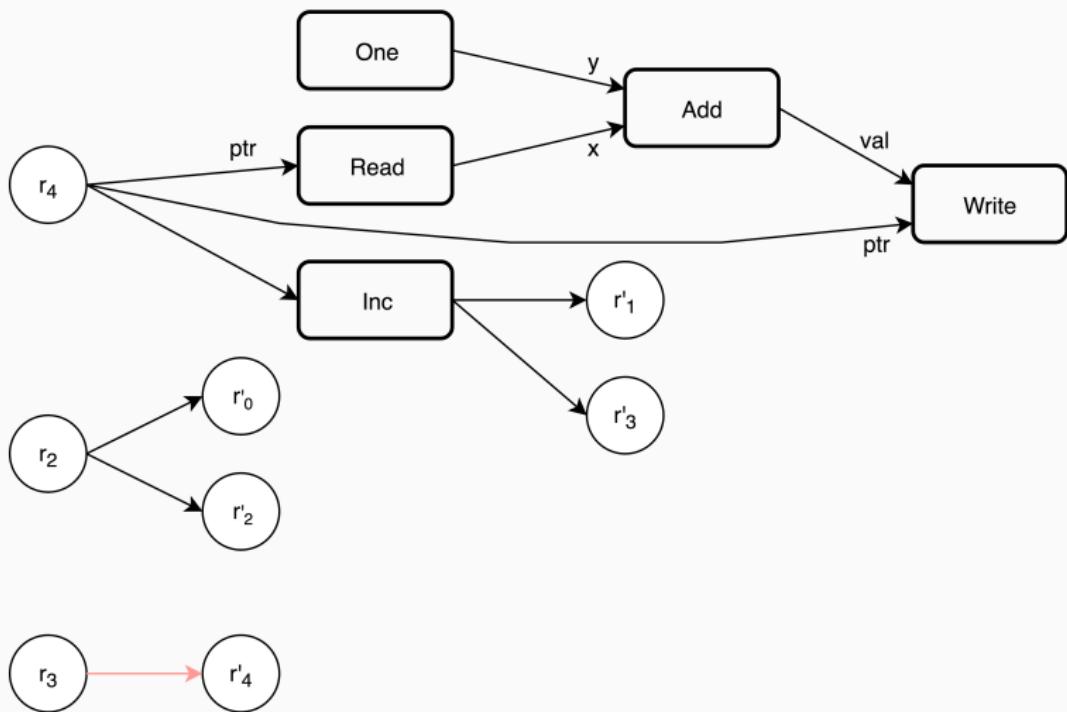
Timestep execution - Circuit execution example



Timestep execution - Circuit execution example



Timestep execution - Circuit execution example



Memory augmented version

Controller

Registers

Modules

Read

Zero

One

Two

Inc

Add

Sub

Dec

Less-
Than

Less-Equal
Than

Equal-
Than

Min

Max

Write

Memory

Memory

The memory is a support of the Neural Random-Access Machines of I memory cells, formalized as $\mathcal{M} \in \mathbb{R}_j^I$. Every value $\mathcal{M}_{i,j}$ is the probability that the i^{th} cell contains the j^{th} integer value of the set N .

Read & Write modules

To interact with the memory, the NRAM has two further modules:

- **Read**: takes a pointer and returns $\mathcal{M}^T p$. In other words, the value(s) pointed by the pointer.

Read & Write modules

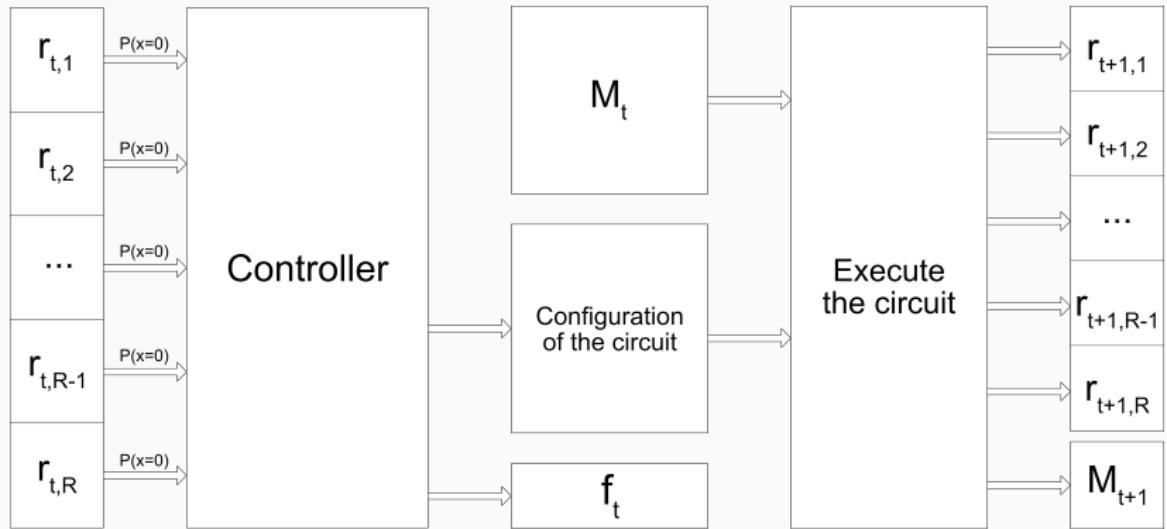
To interact with the memory, the NRAM has two further modules:

- **Read**: takes a pointer and returns $\mathcal{M}^T p$. In other words, the value(s) pointed by the pointer.
- **Write**: takes a pointer and a value, returns zero. Behave as follows:

$$\mathcal{M} = (J - p)J^T \cdot M + pa^T$$

where $J \in 1^l$ and \cdot is the coordinate-wise multiplication.

Timestep execution



Termination of NRAM

The NRAM could terminate its execution in two ways:

- Reaching the last timestep T

Termination of NRAM

The NRAM could terminate its execution in two ways:

- Reaching the last timestep T
- Through its internal system of termination

Internal system of termination

NRAM uses a novel way of termination. In each timestep emits along the circuit configuration an $f_t = \sigma(x_i)$ which represents the willingness of terminating the execution. Although not specified in [Kurach et al., 2015], we set the termination if $f_t \geq 1.0$.

Cost calculation

Starting from f_t , the probability that the execution is not finished in the previous timestep is $\prod_{i=1}^{t-1} (1 - f_i)$, from which are computed

- the probability that the execution is produced in the current timestep

$$p_t = f_t \cdot \prod_{i=1}^{t-1} (1 - f_i)$$

Cost calculation

Starting from f_t , the probability that the execution is not finished in the previous timestep is $\prod_{i=1}^{t-1} (1 - f_i)$, from which are computed

- the probability that the execution is produced in the current timestep

$$p_t = f_t \cdot \prod_{i=1}^{t-1} (1 - f_i)$$

- the probability that the execution is produced in the last timestep

$$p_T = 1 - \sum_{i=1}^{T-1} p_i$$

Cost calculation

Let $\mathcal{M} \in \mathbb{R}^l$ the memory and $\mathcal{EM} \in N^l$, the expected negative log-likelihood is used as cost function and is defined as follows

$$-\sum_{i=1}^T \left(p_t \cdot \sum_{j=1}^{|N|} \log(\mathcal{M}_{i,\mathcal{EM}_i}^{(t)}) \right)$$

Cost calculation - Example

Let $t = 0$, $l = 4$ and $p_0 = 0.5$, we might have

Expected Memory

| | | | |
|---|---|---|---|
| 2 | 0 | 3 | 1 |
|---|---|---|---|

Memory

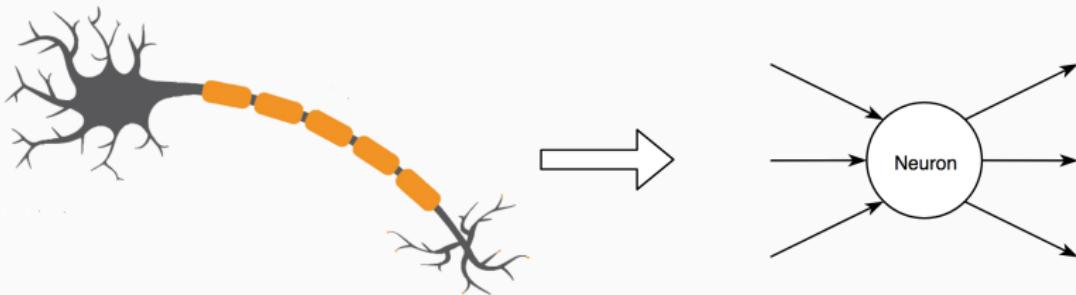
| | | | |
|-----|-----|-----|-----|
| 0.3 | 0.2 | 0.1 | 0.2 |
| 0.7 | 0.1 | 0.0 | 0.2 |
| 0.2 | 0.2 | 0.6 | 0.0 |
| 0.1 | 0.1 | 0.1 | 0.7 |

$$= -(0.5(\log(0.1) + \log(0.7)) + \log(0.0 + \epsilon) + \log(0.1)) + \dots$$

Artificial Neural Network

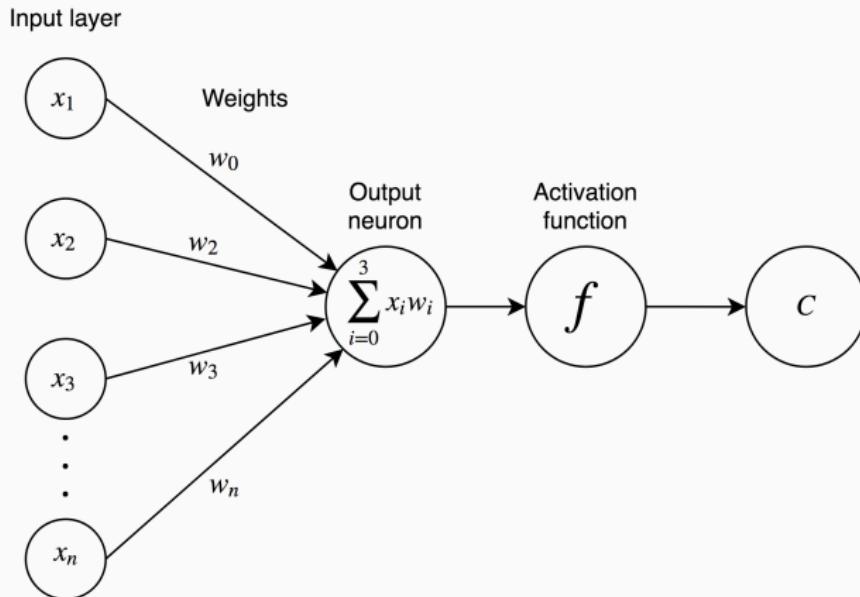
What are?

The Artificial Neural Network are computing systems inspired by biology and belonging to classification techniques set. Classification systems learn patterns by view examples, i.e. approximate a function f that maps example x to the label y , where $(x, y) \in D$.



Perceptron

The most simple and old ANN [Rosenblatt, 1958] is the Perceptron, formed as follows



Perceptron objective

The objective of the Perceptron is to divide linearly the examples approximating the underling function that maps x to y .

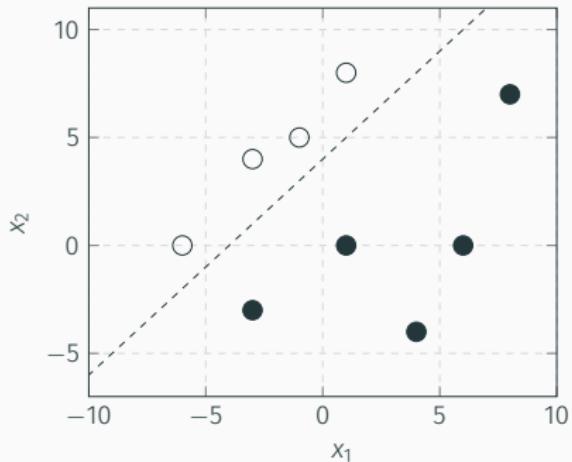
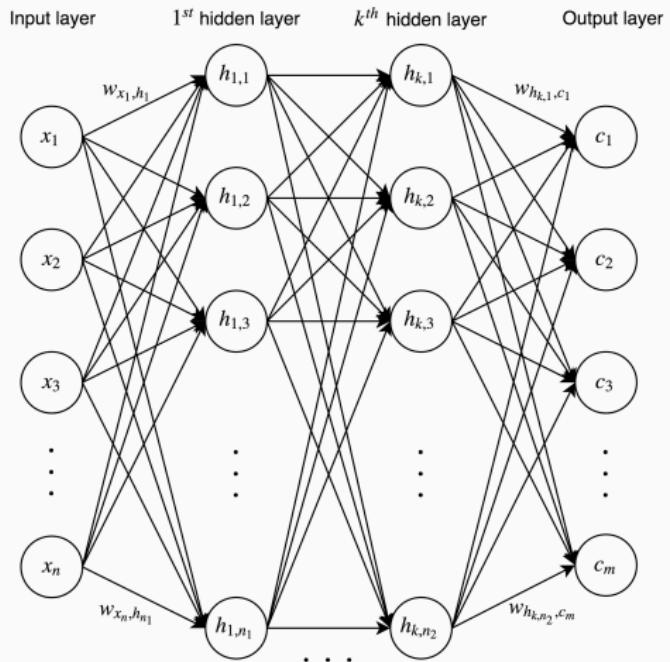


Figure 1: Example with only two features, x_1 and x_2 , and two classes, represented by the black and white circles.

Multi Layer Perceptron (MLP)

An evolution of the Perceptron is the Multi Layer Perceptron



Training of an Artificial Neural Network

The training of an ANN follows two steps:

- The initialization of the weights/parameters

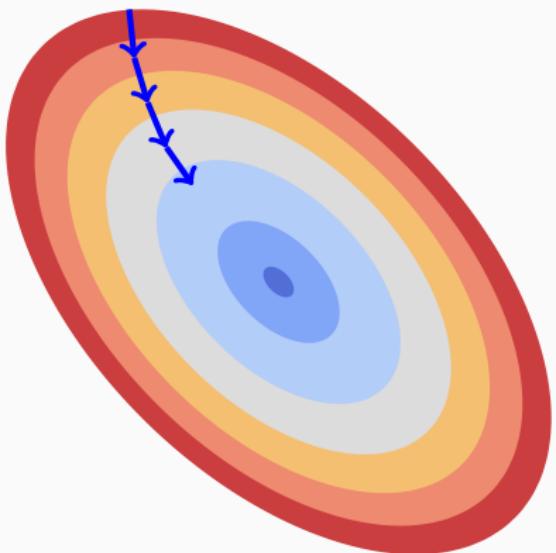
Training of an Artificial Neural Network

The training of an ANN follows two steps:

- The initialization of the weights/parameters
- The re-computation of these searching to minimize the value of an objective function

Gradient Descent & back-propagation

One method to search a optimal/sub-optimal configuration of the parameters through the minimization of the cost function through the Gradient Descent and the back-propagation, with a behaviour similar to the descent of a bowl by an hypothetical ball, where the bowl bottom corresponds to the global minimum of the function.



Back-propagation

The search of those minima is done firstly through the computing of the ANN gradient through the back-propagation, which it is divided into two phases:

- Forward propagation

Back-propagation

The search of those minima is done firstly through the computing of the ANN gradient through the back-propagation, which it is divided into two phases:

- Forward propagation
- Backward propagation

Forward phase

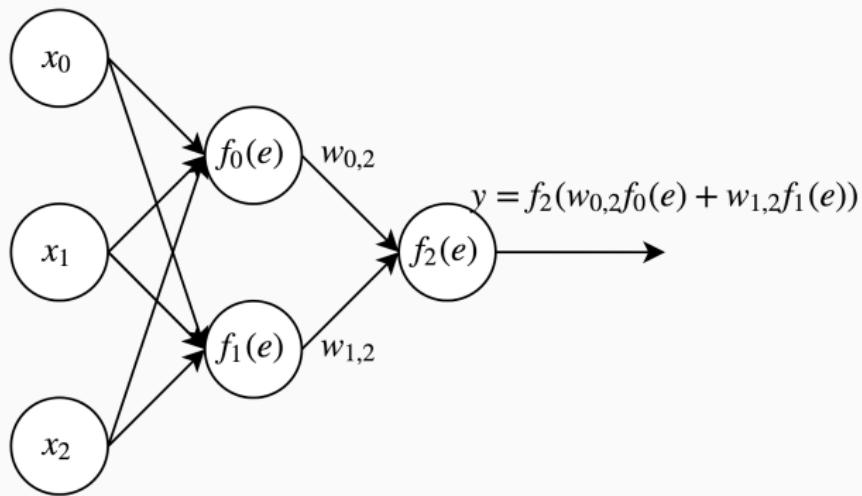


Figure 2: Last pass of forward phase, when the output of ANN is computed

Backpropagation phase

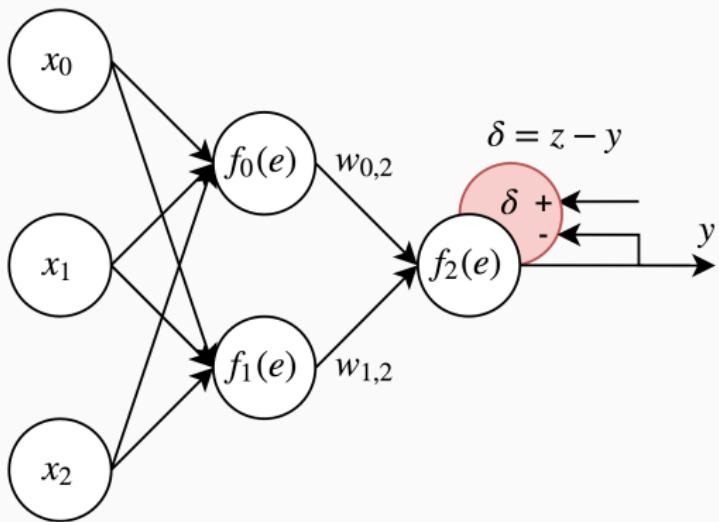


Figure 3: Computing of δ . It can be computed directly because is available a class that can be compared to the ANN exit.

Backpropagation phase

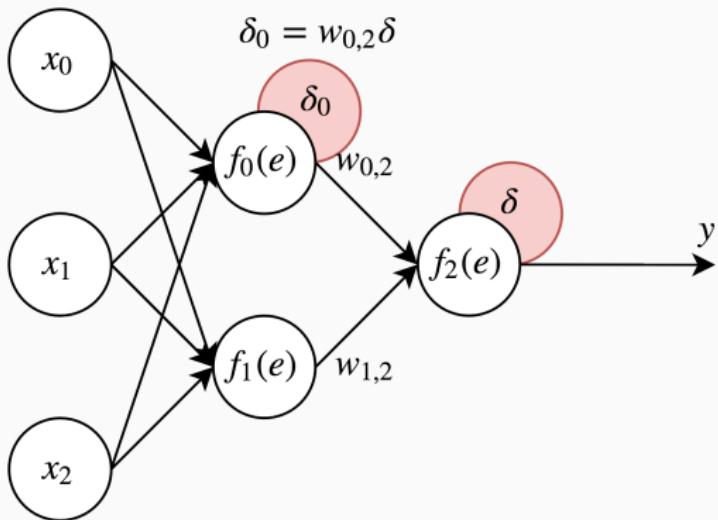


Figure 4: Computing of δ_0 using δ .

Backpropagation phase

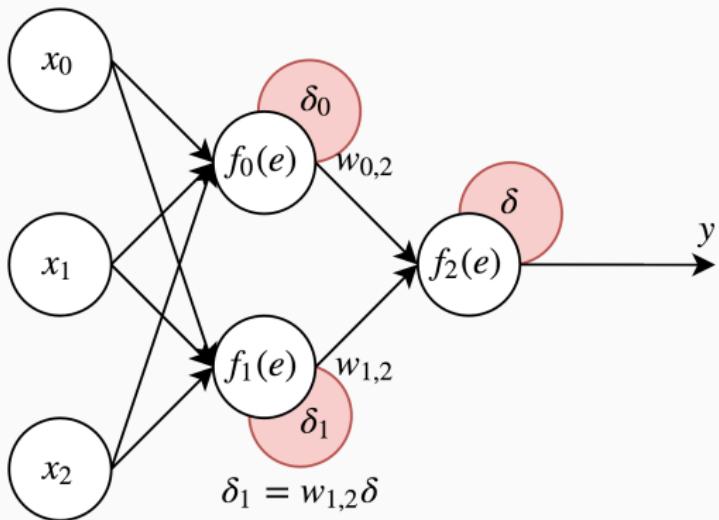


Figure 5: Computing of δ_1 using δ .

Gradient Descent & ADAM

ADAM is a gradient-based optimization algorithm that uses the concept of momentum to regularize the descent of the gradient.

Let β_1 and β_2 two exponential decay rates used for the momentum estimates, ϵ a small real-valued scalar, λ the learning rate and w_{ij} the reference weight. The steps to update a parameter are the following

- Compute $m_{ij}^{(t)}$: $m_{ij}^{(t)} = \beta_1 m_{ij}^{(t-1)} + (1 - \beta_1) \delta_{ij}^{(t)}$

Let β_1 and β_2 two exponential decay rates used for the momentum estimates, ϵ a small real-valued scalar, λ the learning rate and w_{ij} the reference weight. The steps to update a parameter are the following

- Compute $m_{ij}^{(t)}$: $m_{ij}^{(t)} = \beta_1 m_{ij}^{(t-1)} + (1 - \beta_1) \delta_{ij}^{(t)}$
- Compute $v_{ij}^{(t)}$: $v_{ij}^{(t)} = \beta_2 v_{ij}^{(t-1)} + (1 - \beta_2) \delta_{ij}^{2,(t)}$

ADAM

Let β_1 and β_2 two exponential decay rates used for the momentum estimates, ϵ a small real-valued scalar, λ the learning rate and w_{ij} the reference weight. The steps to update a parameter are the following

- Compute $m_{ij}^{(t)}$: $m_{ij}^{(t)} = \beta_1 m_{ij}^{(t-1)} + (1 - \beta_1) \delta_{ij}^{(t)}$
- Compute $v_{ij}^{(t)}$: $v_{ij}^{(t)} = \beta_2 v_{ij}^{(t-1)} + (1 - \beta_2) \delta_{ij}^{2,(t)}$
- Regularize $m_{ij}^{(t)}$ w.r.t. β_1 : $\hat{m}_{ij}^{(t)} = \frac{m_{ij}^{(t)}}{1-\beta_1}$

Let β_1 and β_2 two exponential decay rates used for the momentum estimates, ϵ a small real-valued scalar, λ the learning rate and w_{ij} the reference weight. The steps to update a parameter are the following

- Compute $m_{ij}^{(t)}$: $m_{ij}^{(t)} = \beta_1 m_{ij}^{(t-1)} + (1 - \beta_1) \delta_{ij}^{(t)}$
- Compute $v_{ij}^{(t)}$: $v_{ij}^{(t)} = \beta_2 v_{ij}^{(t-1)} + (1 - \beta_2) \delta_{ij}^{2,(t)}$
- Regularize $m_{ij}^{(t)}$ w.r.t. β_1 : $\hat{m}_{ij}^{(t)} = \frac{m_{ij}^{(t)}}{1-\beta_1}$
- Regularize $v_{ij}^{(t)}$ w.r.t. β_2 : $\hat{v}_{ij}^{(t)} = \frac{v_{ij}^{(t)}}{1-\beta_2}$

ADAM

Let β_1 and β_2 two exponential decay rates used for the momentum estimates, ϵ a small real-valued scalar, λ the learning rate and w_{ij} the reference weight. The steps to update a parameter are the following

- Compute $m_{ij}^{(t)}$: $m_{ij}^{(t)} = \beta_1 m_{ij}^{(t-1)} + (1 - \beta_1) \delta_{ij}^{(t)}$
- Compute $v_{ij}^{(t)}$: $v_{ij}^{(t)} = \beta_2 v_{ij}^{(t-1)} + (1 - \beta_2) \delta_{ij}^{2,(t)}$
- Regularize $m_{ij}^{(t)}$ w.r.t. β_1 : $\hat{m}_{ij}^{(t)} = \frac{m_{ij}^{(t)}}{1-\beta_1}$
- Regularize $v_{ij}^{(t)}$ w.r.t. β_2 : $\hat{v}_{ij}^{(t)} = \frac{v_{ij}^{(t)}}{1-\beta_2}$
- Update $w_{ij}^{(t)}$: $w_{ij}^{(t)} = w_{ij}^{(t)} - \lambda \frac{\hat{m}_{ij}^{(t)}}{\sqrt{\hat{v}_{ij}^{(t)}} + \epsilon}$

Differential Evolution & DENN

Description

Differential Evolution (DE) is a metaheuristic introduced in [Storn and Price, 1997], belonging to the family of Evolutionary Algorithms (EAs), that has as objective the searching of a solution through the parallel evolution of a set of candidate solutions.

This set is defined as **population** and it is composed by NP D-dimensional numerical vectors where each of them is called individual, chromosome or genoma.

Description

Its functioning is iterative and after a first phase of initialization of the individuals the following step is executed

- Mutation

Description

Its functioning is iterative and after a first phase of initialization of the individuals the following step is executed

- Mutation
- Crossover

Description

Its functioning is iterative and after a first phase of initialization of the individuals the following step is executed

- Mutation
- Crossover
- Selection

Mutation

Through the mutation phase a new population, called donor set, is created. Let x an individual at the generation G , also called target. A donor is generated combining x with some other individuals through a differential mutation operator.

Mutation - Example

For example, the method **rand/1** work as follows

$$v_{i,G+1} = x_{r_1,G} + F \cdot (x_{r_2,G} - x_{r_3,G})$$

where $r_1, r_2, r_3 \in \{1, 2, \dots, NP\}$ are mutually exclusive indices
and F is a real-valued constant defined by the user.

Crossover

The Crossover phase does nothing else than mixing up one-by-one the donor vectors components with the target vectors with some crossover methods, generating the **trials** vectors set and enhancing the potential diversity of the population (a better exploration).

Crossover - Example

For example, let the target vector $x_{i,G}$ and the mutant $v_{i,G}$ the **bin** strategy work as follows

$$u_{ji,G} = \begin{cases} v_{ji,G}, & \text{if } (\text{randb}(j) \leq CR) \text{ or } j = rnbr(i) \\ x_{ji,G}, & \text{if } (\text{randb}(j) > CR) \text{ and } j \neq rnbr(i) \end{cases} \quad i = 1, 2, \dots, D \quad (4)$$

where $\text{randb}(j)$ is a function which generate a real-valued number for the j^{th} parameter according to binomial distribution, $CR \in [0, 1]$ is the global user-defined crossover constant used as a threshold and $rnbr(i)$ is a function that generate a random index which ensures that is selected at least one parameter of the mutant v_i .

Selection

Each trial is compared one-by-one to the corresponding target using the fitness function - if the target have a smaller cost with respect to the trial, than it is retained as individual of the population of the next generation and vice-versa.

Differential Evolution variants

There exist various DE variants; some of these that are used in the thesis are

- **JADE**: generates F_i and CR_i for each target x_i , with an adaptation system based on two memories denoted as S_F and S_{CR} re-generated for each \mathbf{G} with the best trials;

Differential Evolution variants

There exist various DE variants; some of these that are used in the thesis are

- **JADE**: generates F_i and CR_i for each target x_i , with an adaptation system based on two memories denoted as S_F and S_{CR} re-generated for each \mathbf{G} with the best trials;
- **SHADE**: generates F_i and CR_i for each target x_i , with an adaptation system based on two memories M_F and M_{CR} of *finite* size H , called *historical memories*, updated with S_F and S_{CR} ;

Differential Evolution variants

There exist various DE variants; some of these that are used in the thesis are

- **JADE**: generates F_i and CR_i for each target x_i , with an adaptation system based on two memories denoted as S_F and S_{CR} re-generated for each \mathbf{G} with the best trials;
- **SHADE**: generates F_i and CR_i for each target x_i , with an adaptation system based on two memories M_F and M_{CR} of *finite* size H , called *historical memories*, updated with S_F and S_{CR} ;
- **L-SHADE**: similar to SHADE, but at every generation it reduces linearly the population.

Mutation variants

As for the DE variants, also for the mutation methods exist variants; some used in the thesis are the following

- **DEGL**: generates a donor $u_i = w \cdot \mathbf{G}_{i,G} + (1 - w) \cdot \mathbf{L}_{i,G}$ where:

Mutation variants

As for the DE variants, also for the mutation methods exist variants; some used in the thesis are the following

- **DEGL**: generates a donor $u_i = w \cdot \mathbf{G}_{i,G} + (1 - w) \cdot \mathbf{L}_{i,G}$ where:
 - $w \in (0, 1)$

Mutation variants

As for the DE variants, also for the mutation methods exist variants; some used in the thesis are the following

- **DEGL**: generates a donor $u_i = w \cdot \mathbf{G}_{i,G} + (1 - w) \cdot \mathbf{L}_{i,G}$ where:
 - $w \in (0, 1)$
 - $\mathbf{G}_{i,G} = \mathbf{x}_{i,G} + \alpha \cdot (\mathbf{x}_{g_best,G} - \mathbf{x}_{i,G}) + \beta \cdot (\mathbf{x}_{r_1,G} - \mathbf{x}_{r_2,G})$

Mutation variants

As for the DE variants, also for the mutation methods exist variants; some used in the thesis are the following

- **DEGL**: generates a donor $u_i = w \cdot \mathbf{G}_{i,G} + (1 - w) \cdot \mathbf{L}_{i,G}$ where:
 - $w \in (0, 1)$
 - $\mathbf{G}_{i,G} = \mathbf{x}_{i,G} + \alpha \cdot (\mathbf{x}_{g_best,G} - \mathbf{x}_{i,G}) + \beta \cdot (\mathbf{x}_{r_1,G} - \mathbf{x}_{r_2,G})$
 - $\mathbf{L}_{i,G} = \mathbf{x}_{i,G} + \alpha \cdot (\mathbf{x}_{n_best_i,G} - \mathbf{x}_{i,G}) + \beta \cdot (\mathbf{x}_{p,G} - \mathbf{x}_{q,G})$

Mutation variants

As for the DE variants, also for the mutation methods exist variants; some used in the thesis are the following

- **DEGL:** generates a donor $u_i = w \cdot \mathbf{G}_{i,G} + (1 - w) \cdot \mathbf{L}_{i,G}$ where:
 - $w \in (0, 1)$
 - $\mathbf{G}_{i,G} = \mathbf{x}_{i,G} + \alpha \cdot (\mathbf{x}_{g_best,G} - \mathbf{x}_{i,G}) + \beta \cdot (\mathbf{x}_{r_1,G} - \mathbf{x}_{r_2,G})$
 - $\mathbf{L}_{i,G} = \mathbf{x}_{i,G} + \alpha \cdot (\mathbf{x}_{n_best_i,G} - \mathbf{x}_{i,G}) + \beta \cdot (\mathbf{x}_{p,G} - \mathbf{x}_{q,G})$
- **Current-to-pbest:** exists in two versions:

Mutation variants

As for the DE variants, also for the mutation methods exist variants; some used in the thesis are the following

- **DEGL**: generates a donor $u_i = w \cdot \mathbf{G}_{i,G} + (1 - w) \cdot \mathbf{L}_{i,G}$ where:
 - $w \in (0, 1)$
 - $\mathbf{G}_{i,G} = \mathbf{x}_{i,G} + \alpha \cdot (\mathbf{x}_{g_best,G} - \mathbf{x}_{i,G}) + \beta \cdot (\mathbf{x}_{r_1,G} - \mathbf{x}_{r_2,G})$
 - $\mathbf{L}_{i,G} = \mathbf{x}_{i,G} + \alpha \cdot (\mathbf{x}_{n_best_i,G} - \mathbf{x}_{i,G}) + \beta \cdot (\mathbf{x}_{p,G} - \mathbf{x}_{q,G})$
- **Current-to-pbest**: exists in two versions:
 - without archive:

$$\mathbf{u}_{i,G} = \mathbf{x}_{i,G} + F_i \cdot (\mathbf{x}_{best,G}^p - \mathbf{x}_{i,G}) + F_i \cdot (\mathbf{x}_{r_1,G} - \mathbf{x}_{r_2,G})$$

Mutation variants

As for the DE variants, also for the mutation methods exist variants; some used in the thesis are the following

- **DEGL:** generates a donor $u_i = w \cdot \mathbf{G}_{i,G} + (1 - w) \cdot \mathbf{L}_{i,G}$ where:
 - $w \in (0, 1)$
 - $\mathbf{G}_{i,G} = \mathbf{x}_{i,G} + \alpha \cdot (\mathbf{x}_{g_best,G} - \mathbf{x}_{i,G}) + \beta \cdot (\mathbf{x}_{r_1,G} - \mathbf{x}_{r_2,G})$
 - $\mathbf{L}_{i,G} = \mathbf{x}_{i,G} + \alpha \cdot (\mathbf{x}_{n_best_i,G} - \mathbf{x}_{i,G}) + \beta \cdot (\mathbf{x}_{p,G} - \mathbf{x}_{q,G})$
- **Current-to-pbest:** exists in two versions:
 - without archive:
$$\mathbf{u}_{i,G} = \mathbf{x}_{i,G} + F_i \cdot (\mathbf{x}_{best,G}^p - \mathbf{x}_{i,G}) + F_i \cdot (\mathbf{x}_{r_1,G} - \mathbf{x}_{r_2,G})$$
 - with an archive A:
$$\mathbf{u}_{i,G} = \mathbf{x}_{i,G} + F_i \cdot (\mathbf{x}_{best,G}^p - \mathbf{x}_{i,G}) + F_i \cdot (\mathbf{x}_{r_1,G} - \tilde{\mathbf{x}}_{r_2,G})$$

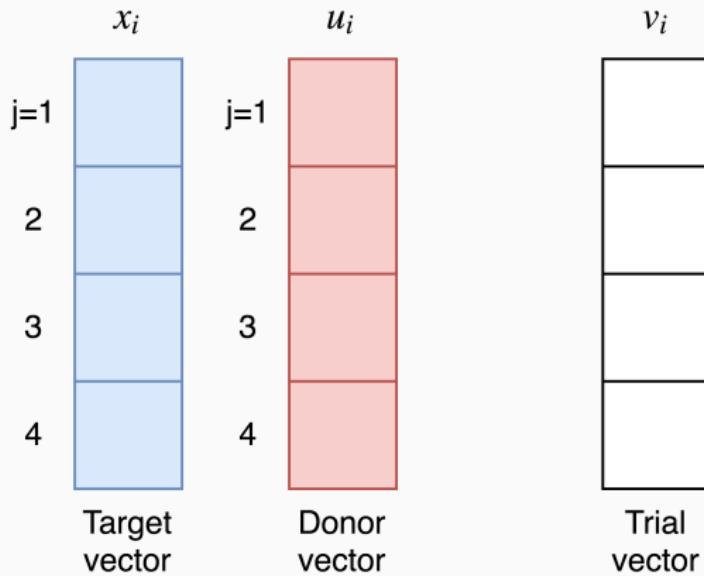
where $\tilde{\mathbf{x}}_{r_2,G}$ is selected in the set $NP \cup A$.

Crossover variant

As for the DE and mutation operators exist more crossover variants, but we used in this case only the method **bin** introduced previously.

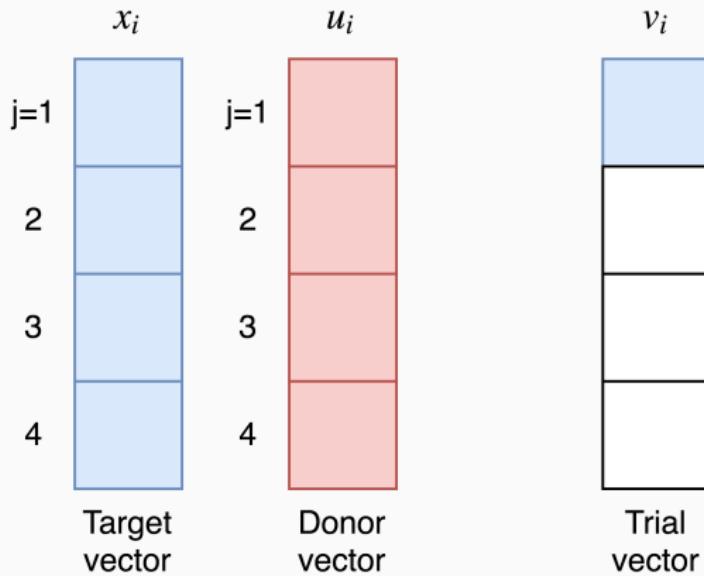
Crossover variant

As for the DE and mutation operators exist more crossover variants, but we used in this case only the method **bin** introduced previously.



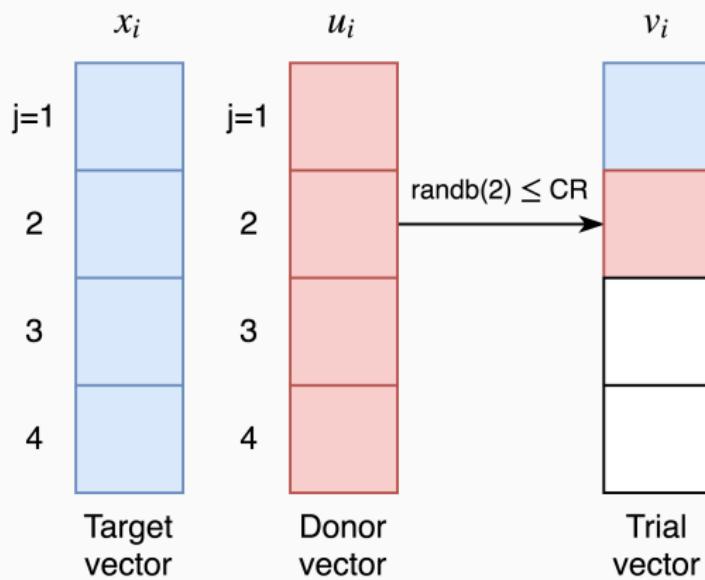
Crossover variant

As for the DE and mutation operators exist more crossover variants, but we used in this case only the method **bin** introduced previously.



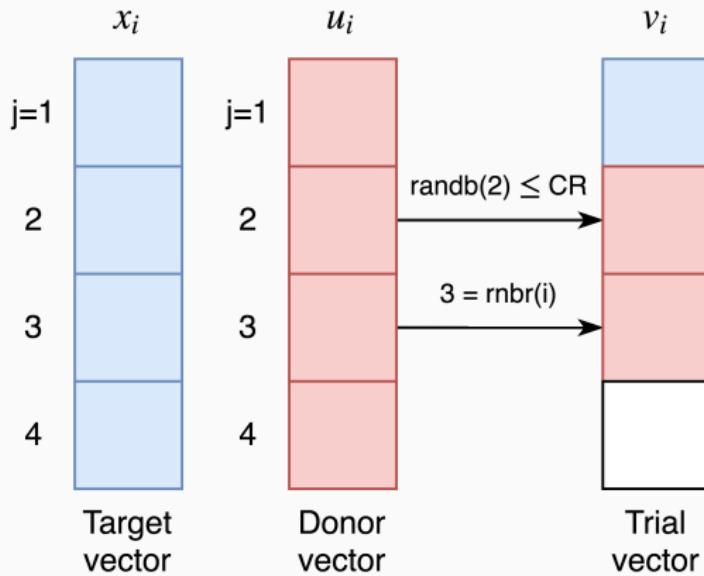
Crossover variant

As for the DE and mutation operators exist more crossover variants, but we used in this case only the method **bin** introduced previously.



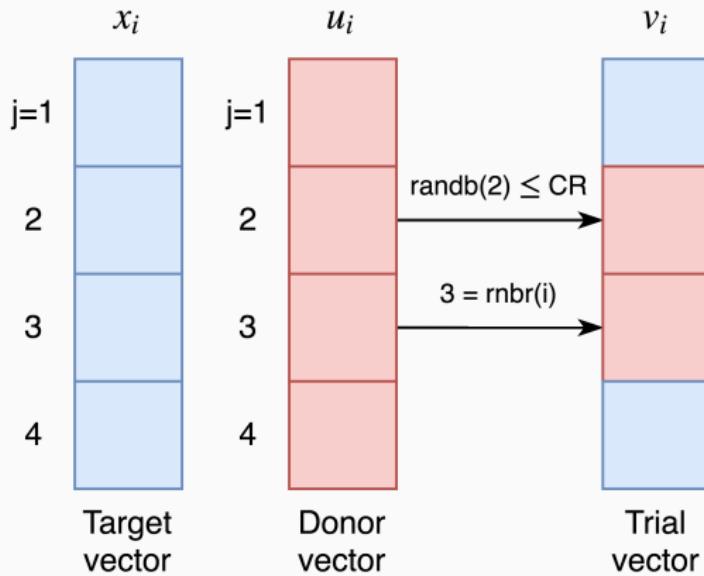
Crossover variant

As for the DE and mutation operators exist more crossover variants, but we used in this case only the method **bin** introduced previously.



Crossover variant

As for the DE and mutation operators exist more crossover variants, but we used in this case only the method **bin** introduced previously.



DENN (Differential Evolution for Neural Network)

DENN is a framework which implements the concepts of the Differential Evolution to train ANN, for now only MLP. Created by Gabriele Di Bari and Mirco Tracolli, it is written in C++ based on the library Eigen. It has a modular structure, permitting so to decide of how the training should be done.

DENN (Differential Evolution for Neural Network)

Since the population is composed by ANN, Differential Evolution is not applied as is but over a set of sets of weights and biases.

For example, let NP the population, **Rand/1** the mutation strategy, r_1, r_2 and r_3 three mutually exclusive indexes, the donor \mathbf{u}_1 is created as follows:

$$\mathbf{u}_1^{w_1} = \mathbf{x}_{r_1}^{w_1} + F \cdot (\mathbf{x}_{r_2}^{w_1} - \mathbf{x}_{r_3}^{w_1})$$

$$\mathbf{u}_1^{b_1} = \mathbf{x}_{r_1}^{b_1} + F \cdot (\mathbf{x}_{r_2}^{b_1} - \mathbf{x}_{r_3}^{b_1})$$

...

$$\mathbf{u}_1^{w_{|\text{HL}|}} = \mathbf{x}_{r_1}^{w_{|\text{HL}|}} + F \cdot (\mathbf{x}_{r_2}^{w_{|\text{HL}|}} - \mathbf{x}_{r_3}^{w_{|\text{HL}|}})$$

$$\mathbf{u}_1^{b_{|\text{HL}|}} = \mathbf{x}_{r_1}^{b_{|\text{HL}|}} + F \cdot (\mathbf{x}_{r_2}^{b_{|\text{HL}|}} - \mathbf{x}_{r_3}^{b_{|\text{HL}|}})$$

Implementation & results

Implementation

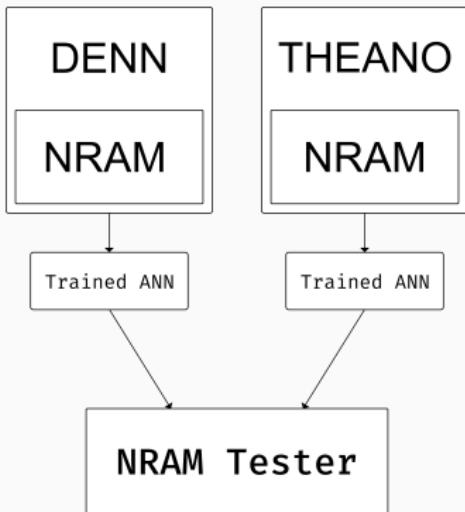


Figure 6: Three implementations were developed: the first two are used to train the controllers and the last is used to test the generalization ability of the discovered controllers.

Datasets

The datasets used during the training are automatically generated at runtime. Each batch of samples is composed by:

- **initial memories:** the memories which are manipulated by NRAM during the execution;
- **expected memories:** the memories which are compared to the modified initial memories for cost computation;
- **cost masks:** used to take into account only parts of the memories during the computing of the cost;
- **error rate masks:** used to take into account only parts of the memories during the computing of the error rate;

Additional used techniques

To help the optimization algorithm searching a solution, we used as in [Kurach et al., 2015] additional technique in the training:

- **Gradient clipping:** used only in the Theano implementation, we clipped the gradient in the interval $[C_1, C_2]$ to avoid the common problem of the Gradient Explosion;

Additional used techniques

To help the optimization algorithm searching a solution, we used as in [Kurach et al., 2015] additional technique in the training:

- **Gradient clipping:** used only in the Theano implementation, we clipped the gradient in the interval $[C_1, C_2]$ to avoid the common problem of the Gradient Explosion;
- **Noise:** used only in the Theano implementation, is added a noise to the computed gradient to enhance the exploration;

Additional used techniques

To help the optimization algorithm searching a solution, we used as in [Kurach et al., 2015] additional technique in the training:

- **Gradient clipping:** used only in the Theano implementation, we clipped the gradient in the interval $[C_1, C_2]$ to avoid the common problem of the Gradient Explosion;
- **Noise:** used only in the Theano implementation, is added a noise to the computed gradient to enhance the exploration;
- **Curriculum Learning:** used only in the DENN implementation due to the Theano structure, it is used to enhance the learning of the ANN.

How the Curriculum Learning works

Curriculum Learning subdivides the training of the ANN in increasing difficulties, where for difficulty is intended the tuple formed by the size of Max-Int (hence, also the memory size), input sequence size and maximum number of allowed timesteps. It starts with the basilar difficulty and is changed when the error rate $1 - \frac{c}{m} < \gamma$ and some generations are executed.

How the Curriculum Learning works

The selection of the new difficulty d is dependent to the current difficulty D and works as follows:

- with probability 10%: pick d randomly from the set of all difficulties according to a uniform distribution;

How the Curriculum Learning works

The selection of the new difficulty d is dependent to the current difficulty D and works as follows:

- with probability 10%: pick d randomly from the set of all difficulties according to a uniform distribution;
- with probability 25%: pick d randomly from the set $[1, D + e]$ according to a uniform distribution, with e is generated from a geometric distribution with a success probability of 0.5;

How the Curriculum Learning works

The selection of the new difficulty d is dependent to the current difficulty D and works as follows:

- with probability 10%: pick d randomly from the set of all difficulties according to a uniform distribution;
- with probability 25%: pick d randomly from the set $[1, D + e]$ according to a uniform distribution, with e is generated from a geometric distribution with a success probability of 0.5;
- with probability 65%: set $d = D + e$ where e is sampled as above

Problems

The following are some of the problems introduced in [Kurach et al., 2015]:

- Access

Problems

The following are some of the problems introduced in [Kurach et al., 2015]:

- Access
- Increment

Problems

The following are some of the problems introduced in [Kurach et al., 2015]:

- Access
- Increment
- Copy

Problems

The following are some of the problems introduced in [Kurach et al., 2015]:

- Access
- Increment
- Copy
- Reverse

Access

Given a value k and an array \mathbf{A} , return $\mathbf{A}[k]$. Input is given as $k, \mathbf{A}[0], \dots, \mathbf{A}[n - 1], \text{NULL}$ and the network should replace the first memory cell with $\mathbf{A}[k]$.

Increment

Given an array A, increment all its elements by 1. Input is given as $A[0], \dots, A[n - 1]$, *NULL* and the expected output is $A[0] + 1, \dots, A[n - 1] + 1$.

| Initial memory |
|---------------------------------------|
| 5 5 9 4 7 8 0 0 0 0 |
| Desired memory |
| 6 6 0 4 8 9 0 0 0 0 |
| Cost mask |
| 1 1 1 1 1 1 1 1 1 1 |
| Error mask |
| 1 1 1 1 1 1 0 0 0 0 |

Copy

Given an array and a pointer to the destination, copy all elements from the array to the given location. Input is given as $p, A[0], \dots, A[n - 1]$ where p points to one element after $A[n - 1]$. The expected output is $A[0], \dots, A[n - 1]$ at positions $p, \dots, p + n - 1$ respectively.

| Initial memory | | | | | | | | | |
|----------------|---|---|---|---|----------|---|---|---|---|
| 5 | 5 | 1 | 4 | 7 | <u>0</u> | 0 | 0 | 0 | 0 |
| Desired memory | | | | | | | | | |
| 5 | 5 | 1 | 4 | 7 | 5 | 1 | 4 | 7 | |
| Cost mask | | | | | | | | | |
| 0 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| Error mask | | | | | | | | | |
| 0 | 0 | 0 | 0 | 0 | 1 | 1 | 1 | 1 | 1 |

Reverse

Given an array and a pointer to the destination, copy all elements from the array in reversed order. Input is given as $p, A[0], \dots, A[n - 1]$ where p points one element after $A[n - 1]$. The expected output is $A[n - 1], \dots, A[0]$ at positions $p, \dots, p + n - 1$.

| Initial memory | | | | | | | | | |
|----------------|---|---|---|---|----------|---|---|---|---|
| 5 | 5 | 1 | 4 | 7 | <u>0</u> | 0 | 0 | 0 | 0 |
| Desired memory | | | | | | | | | |
| 5 | 5 | 1 | 4 | 7 | 7 | 4 | 1 | 5 | |
| Cost mask | | | | | | | | | |
| 0 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| Error mask | | | | | | | | | |
| 0 | 0 | 0 | 0 | 0 | 1 | 1 | 1 | 1 | 1 |

Generalization

Generalization means is the ability of a neural network to recognize patterns in new examples which were never seen before. Remembering how the training is made, here *generalization* means the ability of the controller to generate the “right” circuits also with different parameters of the NRAM.

Objective

The objective of this thesis is to compare the performance over the model NRAM of the Differential Evolution, with and without the help of Curriculum Learning, with respect to Gradient Descent, focusing particularly on the generalization ability of the discovered controllers and the performance of DENN.

Results

| Task | Train complexity | | Reached cost 0 | | Train error | | Generalization | |
|-----------|-----------------------------|-------------------------|----------------|----|-------------|----|----------------|---------|
| | No CL | CL | No CL | CL | No CL | CL | No CL | CL |
| Access | $\text{len}(A) = 8, t = 5$ | $\text{len}(A) \leq 10$ | ✓ | ✓ | 0 | 0 | Perfect | Perfect |
| Increment | $\text{len}(A) = 9, t = 4$ | $\text{len}(A) \leq 10$ | ✓ | ✓ | 0 | 0 | Perfect | Perfect |
| Copy | $\text{len}(A) = 5, t = 11$ | $\text{len}(A) \leq 9$ | ✗ | ✗ | — | — | — | — |
| Reverse | $\text{len}(A) = 4, t = 9$ | $\text{len}(A) \leq 8$ | ✓ | ✓ | 0 | 0 | Perfect | Perfect |

Table 1: Results of the tests with DENN. The train complexity represents the maximum length of integers sequence A used in the training. The train error represents the lowest error rate reached by the trained neural networks. The generalization represents the behaviour of the neural networks with memory sequences longer and more timesteps with respect to those used in training. The evaluation is executed as in [Kurach et al., 2015]. With “—” we indicate the lack of informations due to non-convergence of controller.

Results

| Task | Train complexity | Reached cost 0 |
|-----------|-----------------------------|----------------|
| Access | $\text{len}(A) = 8, t = 5$ | × |
| Increment | $\text{len}(A) = 9, t = 4$ | × |
| Copy | $\text{len}(A) = 5, t = 11$ | × |
| Reverse | $\text{len}(A) = 4, t = 9$ | × |

Table 2: Results of the tests of ADAM.

Results

| Task | JADE/DEGL/1 | | JADE/Current-to-pbest/1 | | SHADE/DEGL/1 | | SHADE/Current-to-pbest/1 | |
|-----------|-------------|----|-------------------------|----|--------------|----|--------------------------|----|
| | No CL | CL | No CL | CL | No CL | CL | No CL | CL |
| Access | ✗ | ✓ | ✓ | ✓ | ✗ | ✗ | ✗ | ✓ |
| Increment | ✗ | ✗ | ✓ | ✗ | ✗ | ✗ | ✗ | ✓ |
| Copy | ✗ | ✗ | ✗ | ✗ | ✗ | ✗ | ✗ | ✗ |
| Reverse | ✗ | ✓ | ✓ | ✗ | ✓ | ✓ | ✗ | ✓ |

Table 3: Convergence table. Although not specified, all the variants are associated with the crossover method **bin**. L-SHADE seems unable to reach a solution in all the tasks, probably due to the linear reducing of the population; due to these results it is not shown in this table. For the other cases, except for the task **Increment** with JADE/Current-to-pbest/1/bin, it can be seen that Curriculum Learning boost the training of the controller.

Charts of convergence - Access

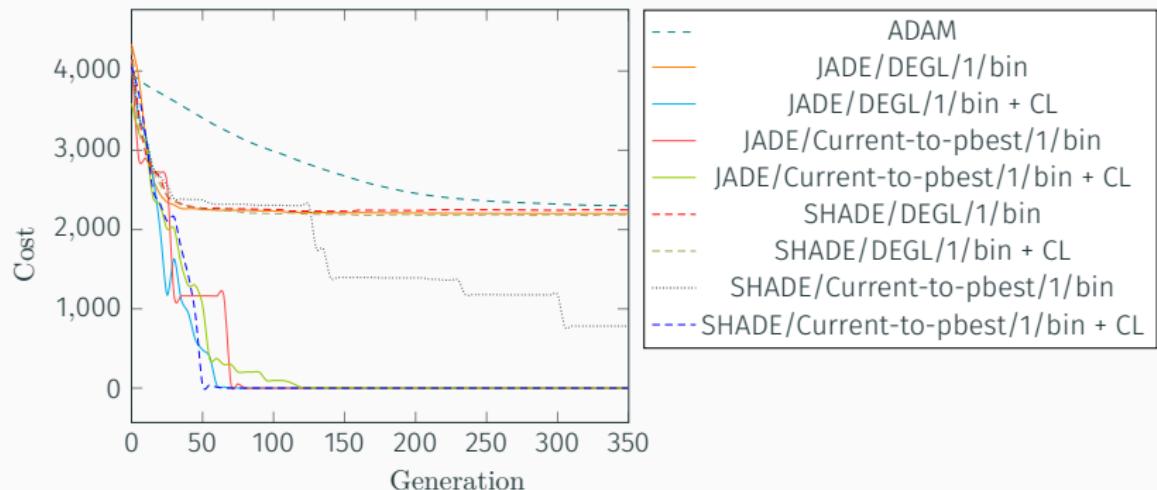


Figure 7: Costs convergence of DE variants with task **Access**.

Charts of convergence - Reverse

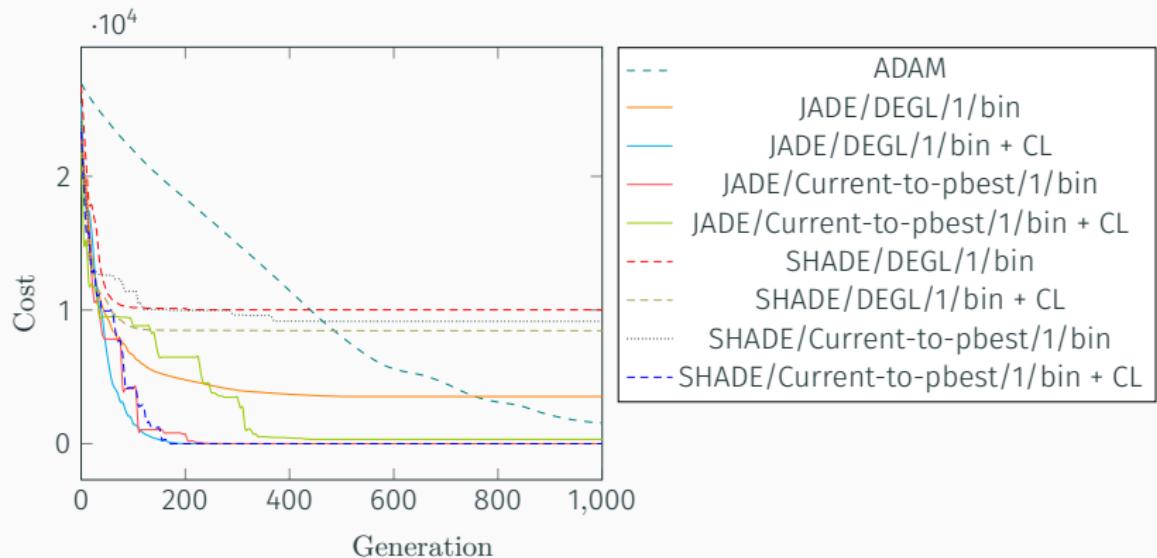
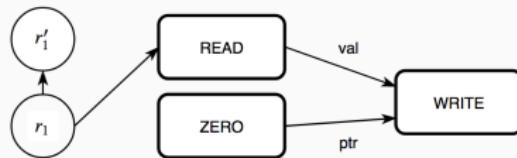


Figure 8: Costs convergence of DE variants with task **Reverse**.

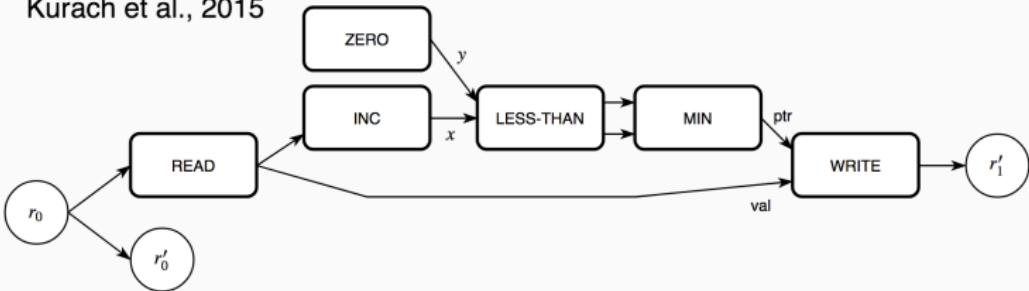
Circuits - Access

DENN with Curriculum Learning



Generated for every timesteps ≥ 2 .

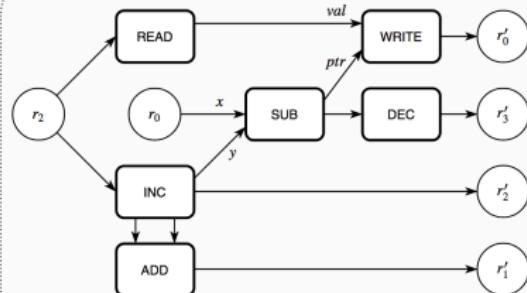
Kurach et al., 2015



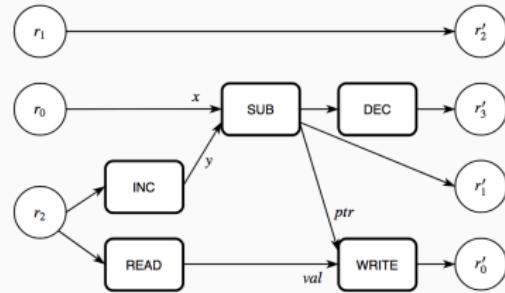
Generated for every timesteps ≥ 2 .

Circuits - Reverse

DENN with Curriculum Learning

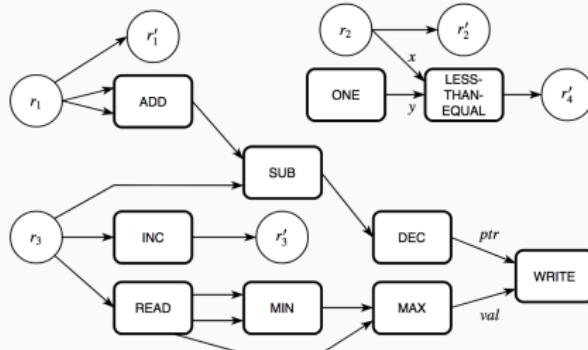


Generated for the $t \in [2, \text{size}(A)]$ and for the odd timesteps $\geq \text{size}(A)$



Generated for the even timesteps $\geq \text{size}(A)$

Kurach et al., 2015



Generated at every timesteps ≥ 2

Thanks for the attention!

References

-  Kurach, K., Andrychowicz, M., and Sutskever, I. (2015).
Neural random-access machines.
CoRR, abs/1511.06392.
-  Rosenblatt, F. (1958).
The perceptron: a probabilistic model for information storage and organization in the brain.
Psychological review, 65 6:386–408.
-  Storn, R. and Price, K. (1997).
Differential evolution - a simple and efficient heuristic for global optimization over continuous spaces.
Journal of Global Optimization, 11.